

Computational Neuroscience

EEE 482/582

Can Kocagil

21602218

Homework-4



Department of Electric & Electronics Engineering

Bilkent University

Ankara, Turkey

25.04.2021

TABLE OF CONTENTS

List of Figures	ii
List of Tables	ii
1. Question 1	1
1.1. Part A	1
1.2. Part B	6
1.3. Part C	13
1.4. Part D	19
2. Question 2	27
2.1. Part A	27
2.2. Part B	29
2.3. Part C	31
2.4. Part D	35
2.5. Part E	37
3. Source Code	40
References	56

LIST OF FIGURES

1	Sample stimuli face images	3
2	Explained Variance Along PCs	4
3	Visualization of First 25Pcs	5
4	Visualization of Original First 36 Images	7
5	Reconstructed 36 images based on first 10 PCs	8
6	Reconstructed 36 images based on first 25 PCs	9
7	Reconstructed 36 images based on first 50 PCs	10
8	Reconstructed 36 images based on first 100 PCs	11
9	Visualization of 10 ICs	14
10	Visualization of 25 ICs	14
11	Visualization of 50 ICs	15
12	FastICA reconstruction based on 10 independent components.....	16
13	FastICA reconstruction based on 25 independent components.....	17
14	FastICA reconstruction based on 50 independent components.....	18
15	Visualization of NNMF 10 MFs	20
16	Visualization of NNMF 25 MFs	21
17	Visualization of NNMF 50 MFs	22
18	NNMF Reconstruction of faces based on 10 MFs	23
19	NNMF Reconstruction of faces based on 25 MFs	24
20	NNMF Reconstruction of faces based on 50MFs	25
21	Gaussian Tuning Curves of a Population of Neurons.....	28
22	Neural Population Response to the Stimulus $x = -1$	29
23	Actual and WTA Estimated Stimuli Across Trials	31
24	Actual and MLE Estimated Stimuli Across Trials	34
25	Actual and MAP Estimated Stimuli Across Trials	36
26	Histogram and Density Plot of Error Rate w.r.t. varying σ	38
27	Plot of Error Rate w.r.t. varying σ	39

LIST OF TABLES

1	Error Statistics based on σ values with MLE	38
---	--	----

1. QUESTION 1

In this question, stimuli consisting of face images have been used in a study on visual perception and experimental stimuli are provided in the file `hw4_data1.mat`, which contains face images down-sampled to a 32 x 32 square grid.

1.1. Part A. In part a, the researcher would like to fit encoding models between the stimuli (i.e., face images) and the measured neural responses. To do that, one first needs to define explanatory variables (i.e., regressors) that capture important variations in stimulus properties during the experiments. To accomplish this goal, we will perform Principal Component Analysis (PCA) on face images. Then, we will explore the the captured variance by PCA with different number of principal components.

The curse of dimensionality is a common problem for ML pipelines and refer to various phenomena that arise when analyzing and organizing data in high-dimensional spaces that do not occur in low-dimensional settings [1]. Moreover, high dimensional data manifest during analyzing or visualizing the data to identify patterns, and some manifest while training machine learning models. Hence, many ML algorithms are prone to higher dimensional feature spaces so that they are likely to overfit in a statistical sense. Simultaneously, when the researcher would like to fit encoding models between the stimuli and the measured neural responses, the curse of dimensionality is also neuroscientifically challenged problem. Hence, before encoding the process, eliminating some features or projecting them to lower dimensional space is critical. In this work, we'll work with PCA algorithm to reduce the dimensionality of the face images while trying to keep the information in the original feature space.

PCA is an unsupervised, non-parametric machine learning algorithm primarily used for dimensionality reduction. By PCA, we can extract information from high dimensional features space by projecting it into lower dimensional subspace while preserving the most of the variation in the data. Hence, the aim of the PCA algorithm is to produce principal components (PCs) that are orthogonal in direction that capture the most of the variance in the given data. It is linear algorithm seeks for the greatest variability in the orthogonal (uncorrelated) directions in the subspace. Hence, the underlying assumption in PCA is that data lies on or near a low d-dimensional linear subspace.

Let's deep dive into PCA. Principal Components (PC) are orthogonal directions that capture most of the variance in the data. For example, 1st PC represent the direction of greatest variability in data. In other means, projection of data points along 1st PC discriminate the data most along any one direction. Let v_1, \dots, v_d denote the principal directions such that they are orthogonal and unit norm.

$$(1) \quad v_i^T v_j = 0 \text{ and } v_i^T v_i = 1 \quad i \neq j \quad \forall i, j$$

Then, the ultimate aim is to find a vector that maximizes the sample variance of the projections.

$$(2) \quad \max_v \frac{1}{n} \sum_{i=1}^n (v^T x_i)^2 = \max_v v^T X X^T V \text{ s.t. } v^T v = 1$$

Then, the problem becomes constrained optimization problem in the sense of Lagrangian.

$$(3) \quad \text{Lagrangian} : \max_v v^T X X^T v - \lambda v^T v$$

Hence, $\frac{\partial(v^T X X^T v - \lambda v^T v)}{\partial v}$ gives the optimal principal axes by wrapping constraint into the objective function.

$$(4) \quad (X X^T - \lambda I)v = 0 \Leftrightarrow (X X^T)v = \lambda v$$

Then, the problem becomes eigenvalue problem. In these formulation, v becomes the eigenvectors of sample covariance matrix $X X^T$ assuming the data is centered. Hence, the eigenvalue λ denote the amount of variability captured along that dimension. Further, since $\lambda_1 > \lambda_2 > \dots$, the first PC captures the greatest variance in the data and it is associated with the largest eigenvalue. Hence, solving the above eqn (4) as a eigenvalue formulation, gives the principal directions so that we can then project our data into computed principal directions. As another interpretation of PCA, it finds the vectors v such that projections on to the vectors solves the sum of squared, specifically MSE, reconstruction. These reconstructions represents the lower rank estimations of original data space. Finally, once we got principal axes, we project our original data space into these that are called transformed representation projections.

$$(5) \quad \text{Transformed Representation Projections} = [v_1^T x_i, \dots, v_d^T x_i]$$

We're ready to apply PCA to face images. In the question settings, we need to perform PCA on 1000 face images with 100PCs. Before that, let's visualize some sample images to explore faces. Here I also put the libraries that I used along that homework.

```

1 # Imports:
2 import numpy as np, matplotlib.pyplot as plt, scipy.stats as stats, pandas as pd
3 from sklearn.decomposition import PCA, FastICA, NMF
4 import random, h5py
5
6 settings = np.seterr(all='ignore')

```

Hence, I utilize the **scikit-learn** library to perform unsupervised algorithms along that homework. Let's visualize some images and their spatial structures.

```

1 # Retrieving data:
2 faces = h5py.File('hw4_data1.mat','r')['faces'][:, :]
3
4 # Little bit of dimension manipulation for representing images:
5 N, num_pixel = faces.shape
6 image_faces = faces.reshape(N, np.int(np.sqrt(num_pixel)),
    ↳ np.int(np.sqrt(num_pixel)))
7 print(image_faces.shape)
8
9
10 # Let's look at the face images:
11 fig, axs = plt.subplots(3,3, figsize=(8,8))
12 for i, axes in enumerate(axs.flatten()):
13     axes.imshow(image_faces[i], cmap='gray')
14     axes.axis('off')

```



FIGURE 1. Sample stimuli face images

As we can see, spatial resolution of stimuli is low as it reduces the computation complexity of the problem. Anyway, let's move on the PCA computation and variance graph w.r.t. PCs.

```

1 # Latent representation dimension:
2 latent_dim = 100
3 pca = PCA(n_components = latent_dim)
4 principalComponents = pca.fit_transform(faces)
5
6 num2str = lambda x : str(round(sum(x),3))
7
8 legends = [
9
10    pca.explained_variance_ratio_[:10],
11    pca.explained_variance_ratio_[:25],
12    pca.explained_variance_ratio_[:50],
13    pca.explained_variance_ratio_[:]
14
15 ]
16
17 legends = list(map(num2str,legends))
18
  
```

```

19 plt.figure(figsize = (10,5))
20 plt.plot(pca.explained_variance_ratio_, color = 'r')
21 plt.xlabel('PCs')
22 plt.ylabel('Explained Variance')
23 title = 'Principal Components versus Explained Variance \n'
24 title += 'First (10,25,50,100) PCs explained variance = '
25 title += ' (' + ' '.join(legends) + ')'
26 plt.title(title)
27 plt.grid()
28 plt.show()
29
30 pca_logs = f'Variance explained by PCA for 10 components {legends[0]}\n'
31 pca_logs += f'Variance explained by PCA for 25 components {legends[1]}\n'
32 pca_logs += f'Variance explained by PCA for 50 components {legends[2]}\n'
33 pca_logs += f'Variance explained by PCA for 100 components {legends[3]}'
34 print(pca_logs)

```

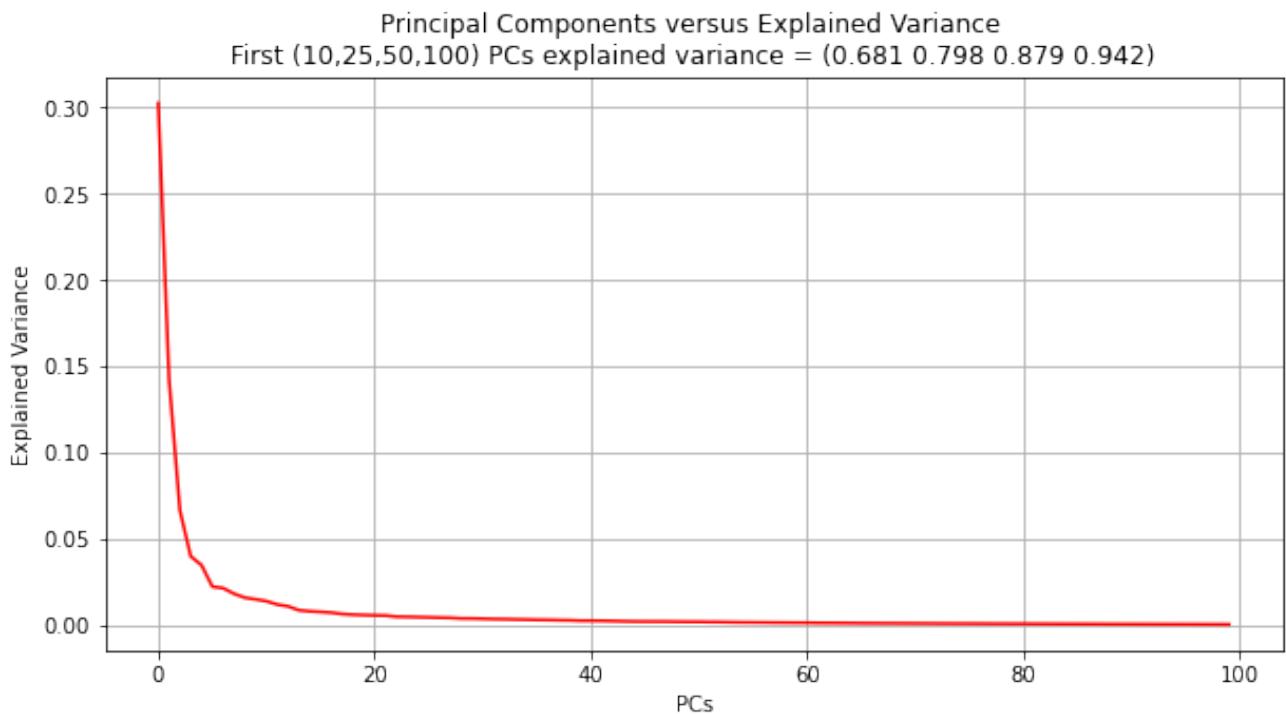


FIGURE 2. Explained Variance Along PCs

Variance explained by PCA for 10 components 0.681
 Variance explained by PCA for 25 components 0.798
 Variance explained by PCA for 50 components 0.879
 Variance explained by PCA for 100 components 0.942

As we can see, with the increased number of PCs, we capture greatest number of variance in the data. With just a first 10 PCs, we can explain more than 68% variance in the data. Further, with the first 100 PCs, we can almost capture the variance of the original data space as explained more than 94%.

Then, the question asks display the first 25PCs. Let's visualize the first 25PCs as follows.

```
1 fig, axes = plt.subplots(5, 5, figsize=(12,12))
2 for i, ax in enumerate(axes.flat):
3     ax.imshow(pca.components_[i].reshape(32, 32).T, cmap = 'gray')
4     ax.axis('off')
```



FIGURE 3. Visualization of First 25Pcs

We can see that first 25PCs captures spatial structure of face images with representative facial points such as eyes, nose, eyebrow. Moreover, these are called **eigen faces** that refer to set of eigenvectors and used in face & facial recognition applications in the context of computer vision.

1.2. Part B. Then, the researcher would like to know how many PCs are sufficient to obtain a reasonable representation of the stimuli. To explore this, we reconstruct images based on PCs. In other words, we'll estimate the original data space with low rank method based on the computed PCs. Here, the question asks to reconstruct based on first 10, 25, and 50 PCs with their corresponding visualizations. Then, we'll report the error statistics based on MSE loss. Here, I construct 2 utility functions for PCA reconstruction and formatted face plotting as follows.

```

1 def pca_reconstruction(data:np.ndarray,
2                         trained_pca:PCA,
3                         number_PCs:int) -> np.ndarray:
4     """
5         Given the input data, trained PCA variable and # of PCs, reconstruct
6         images based on the PCs components.
7
8         Arguments:
9             - data (np.ndarray) : Input data
10            - trained_pca (PCA) : trained PCA variable
11            - number_PCs (int) : # of PCs to reconstruct images
12
13        Returns:
14            - reconstructed_data (np.ndarray) : Reconstructed/Predicted data via
15                given # of PCs
16        """
17
18    pca_mean = trained_pca.mean_
19    mean_removed = data - pca_mean
20    pca_components = trained_pca.components_[:number_PCs]
21
22    return mean_removed @ pca_components.T @ pca_components + pca_mean
23
24 def plot_faces(faces:np.ndarray,
25                 suptitle:str) -> None:
26     """
27         Given the face matrix and its suptitle, plots the 6x6 grid of faces.
28
29         Arguments:
30             - faces      (np.ndarray) : Face data to be plotted
31             - suptitle   (str)       : Suptitle of the visualization
32
33
34     fig, axes = plt.subplots(6, 6,
35                             figsize=(10,10),
36                             facecolor='white',
37                             subplot_kw= {
38                                 'xticks': [],
39                                 'yticks': []
34     }
35     )
36     fig.suptitle(suptitle,
37                   fontsize = '14')
38
39
40
41
42
43

```

```

44     fig.tight_layout(rect = [0, 0, 1, .95])
45
46     for i, ax in enumerate(axes.flat):
47         ax.imshow(faces[i].reshape(32, 32).T, cmap='gray')
48         ax.set_xlabel(i+1)

```

Hence, we are ready to reconstruct images based on first 10,25 and 50 PCs.

```

1 faces_PCA_10 = pca_reconstruction(faces,pca,10)
2 faces_PCA_25 = pca_reconstruction(faces,pca,25)
3 faces_PCA_50 = pca_reconstruction(faces,pca,50)
4 faces_PCA_100 = pca.inverse_transform(principalComponents)

```

Note that I also compute the reconstruction of face data based on the 100Pcs to explore the further behavior of estimations. From now on, I visualize the reconstructions and original data space with 36 samples with their corresponding comments. Let's start with the original first 36 faces images.

```
1 plot_faces(faces,suptitle = 'Original Versions of the First 36 Images')
```



FIGURE 4. Visualization of Original First 36 Images

```
1 plot_faces(faces_PCA_10,suptitle = 'Reconstructed 36 images based on first 10  
→ PCs')
```

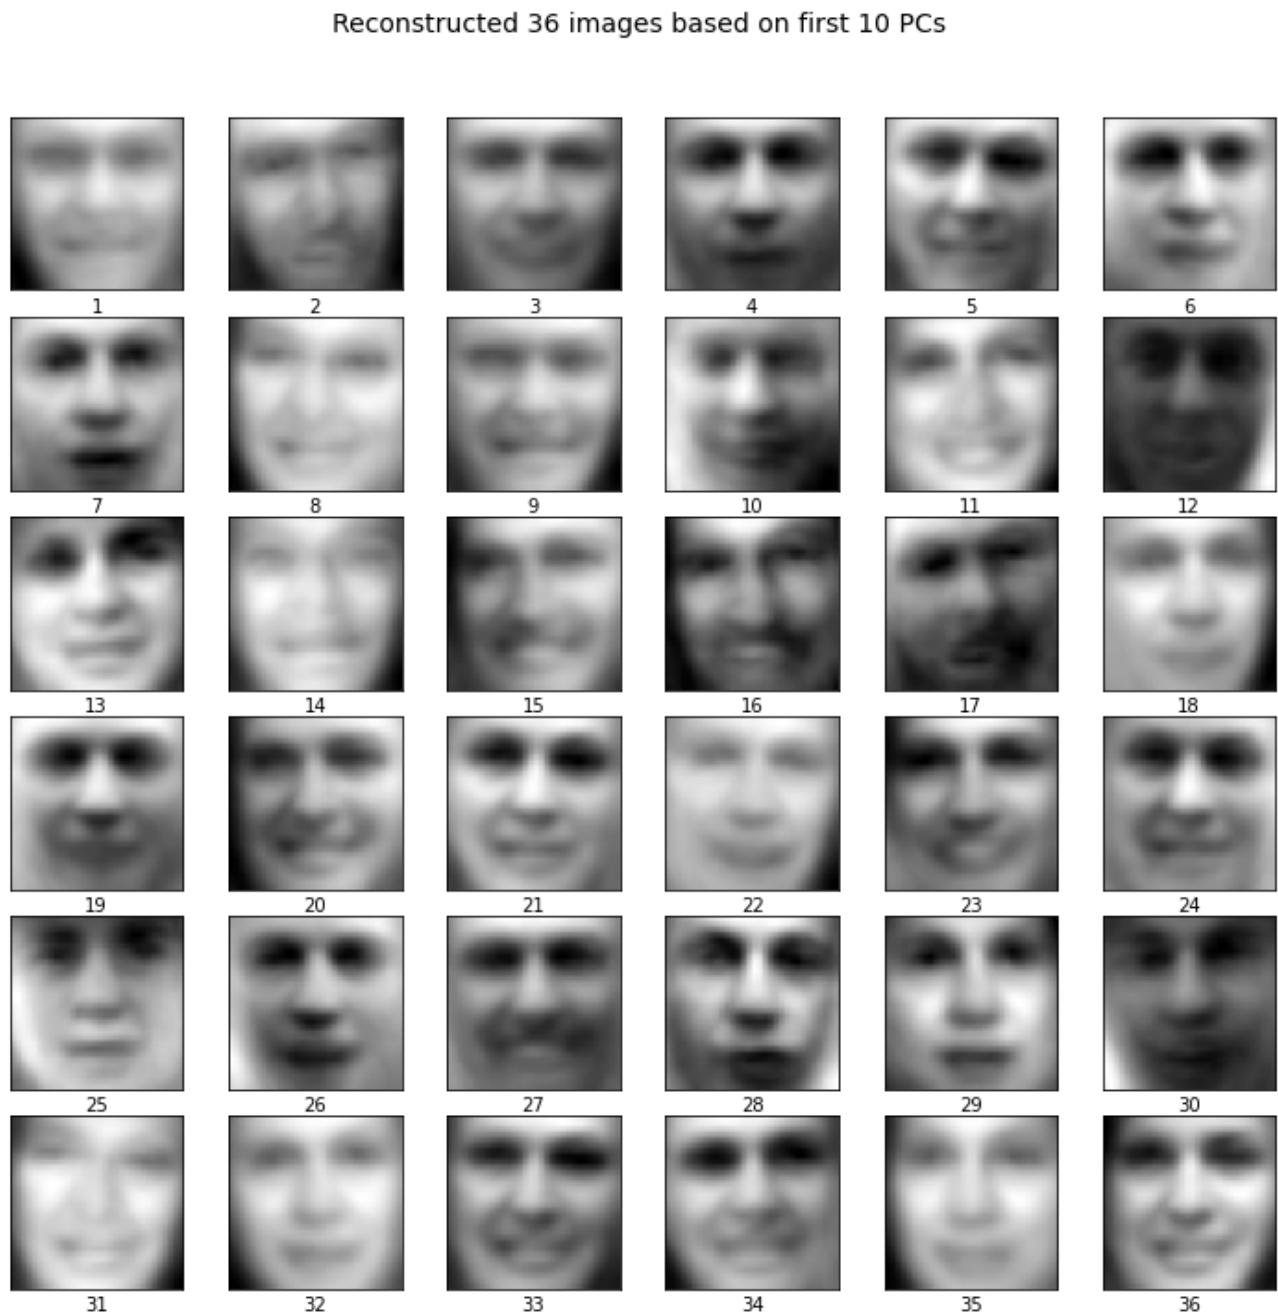


FIGURE 5. Reconstructed 36 images based on first 10 PCs

As we can see, we capture the spatial semantics of the face images just based on the first 10 PCs. The reconstructions are quite enough for early experimental purposes.

```
1 plot_faces(faces_PCA_25,suptitle = 'Reconstructed 36 images based on first 25  
→ PCs')
```

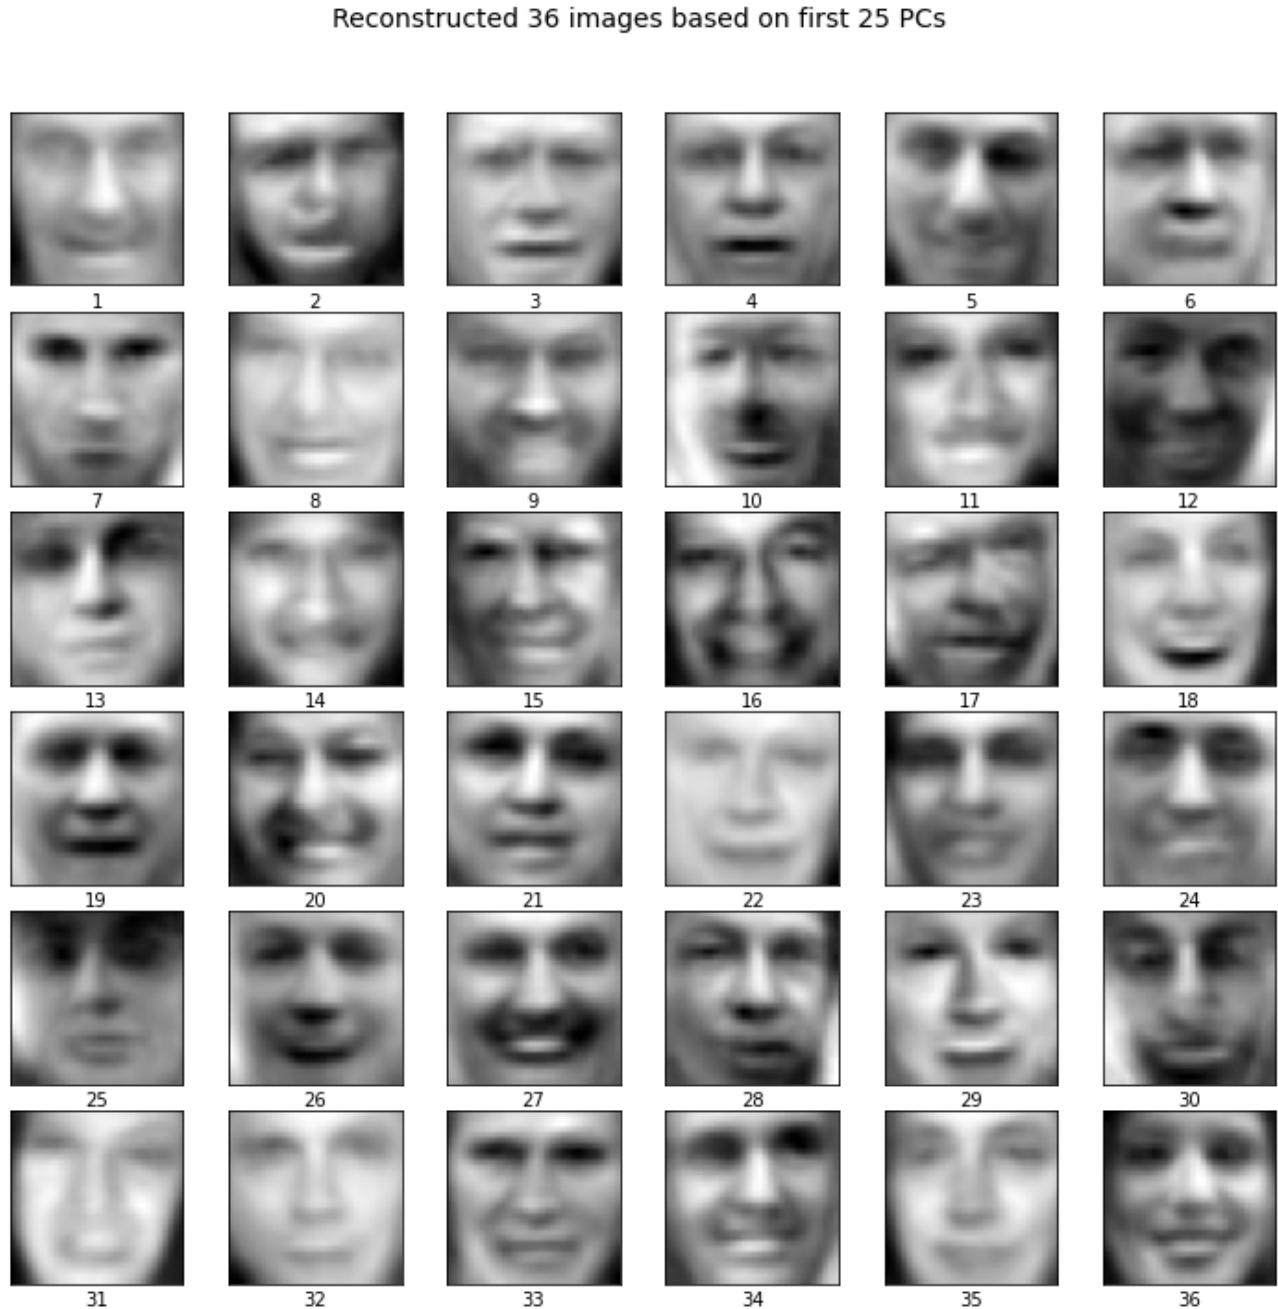


FIGURE 6. Reconstructed 36 images based on first 25 PCs

As we can see, we further capture the spatial semantics of the face images based on the first 25 PCs. The facial structures are much more clear, and we see clear distinctions between mutual faces. Probably, facial matching algorithms will match the low rank estimations and original faces accurately.

```
1 plot_faces(faces_PCA_50,suptitle = 'Reconstructed 36 images based on first 50  
→ PCs')
```

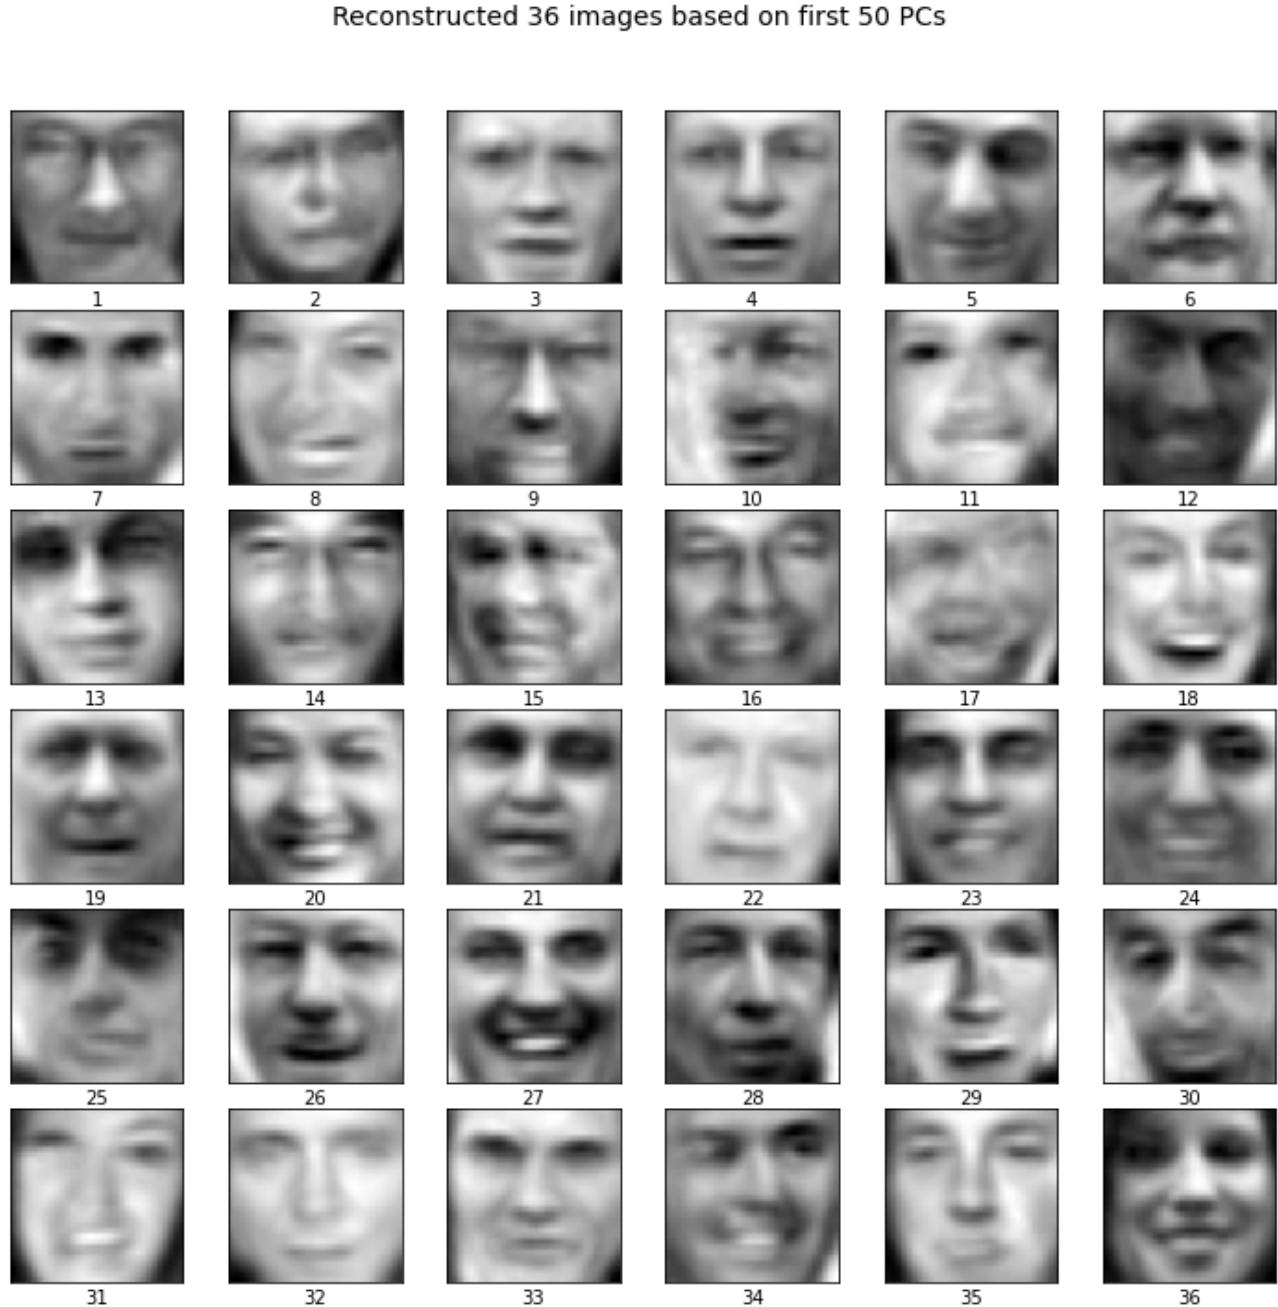


FIGURE 7. Reconstructed 36 images based on first 50 PCs

As we can see, we further capture the high level semantics of the face images based on the first 50 PCs. The spatial structure is clear, faces are identifiable by human vision. Moreover, the facial keypoints are much more clear w.r.t. previous reconstructions. Facial landmark detection algorithms can probably start working with these low rank estimations accurately.

```
1 plot_faces(faces_PCA_100, suptitle = 'Reconstructed 36 images based on first 100  
→ PCs')
```

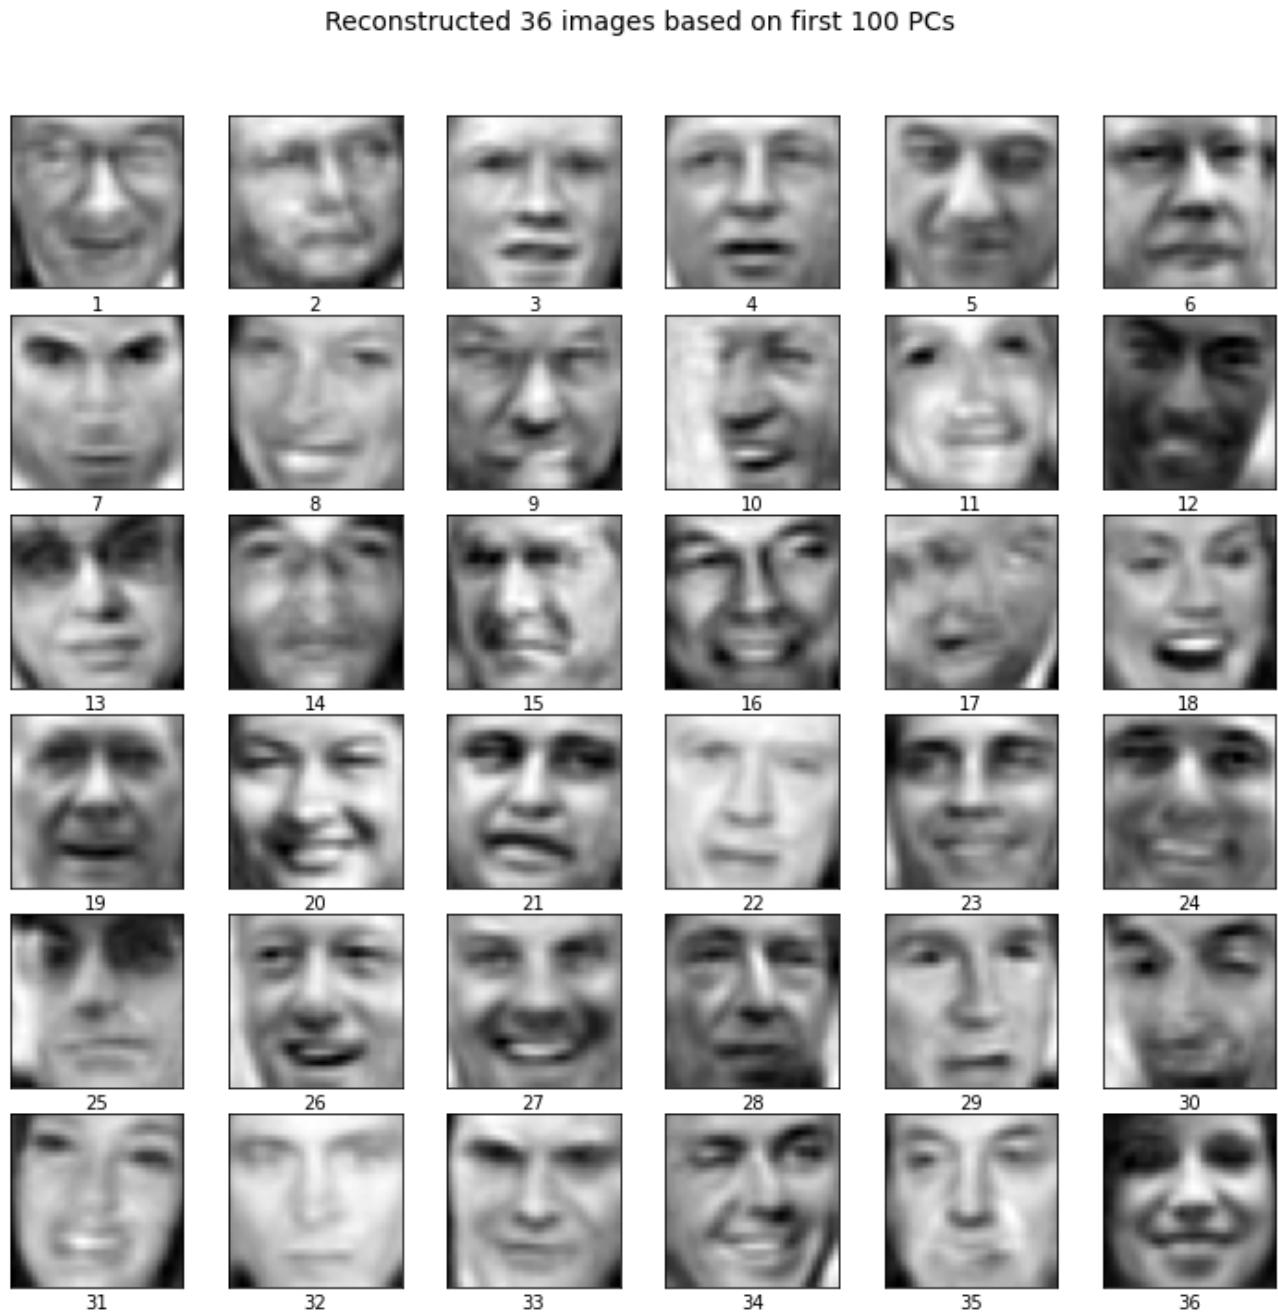


FIGURE 8. Reconstructed 36 images based on first 100 PCs

As we can see, we further capture the semantically meaningful spatial structure of the face images based on the first 100 PCs. The reconstructions are quite well as we capture more than 94% of the variability in the original data space.

Then, let's compute the descriptive statistics of the reconstruction errors based on the MSE loss.

```

1 def squared_error(y_true:np.ndarray,y_pred:np.ndarray) -> np.ndarray:
2     """
3         Given the ground truth matrix and prediction, computes element wise
4         squared error.
5
6     Arguments:
7         - y_true (np.ndarray) : ground truth
8         - y_pred (np.ndarray) : prediction
9
10    Returns:
11        square_error (np.ndarray) : Point-wise MSE loss
12
13    """
14    assert y_true.shape == y_pred.shape, f'Mismatch Dimension!, {y_true.shape}'
15    return (y_true - y_pred) ** 2

```

Then, we compute the average and standard deviation of MSE loss between original face images and their PCA reconstructions as follows.

```

1 mse_10 = squared_error(y_true = faces, y_pred = faces_PCA_10)
2 mse_25 = squared_error(y_true = faces, y_pred = faces_PCA_25)
3 mse_50 = squared_error(y_true = faces, y_pred = faces_PCA_50)
4
5
6 std_mse_10 = mse_10.mean(-1).std()
7 std_mse_25 = mse_25.mean(-1).std()
8 std_mse_50 = mse_50.mean(-1).std()
9
10 mean_mse_10 = mse_10.mean()
11 mean_mse_25 = mse_25.mean()
12 mean_mse_50 = mse_50.mean()
13
14
15 print(f'PCA reconstruction loss stats based on first 10 PCs, \n (mean,std) =
16     {mean_mse_10, std_mse_10}')
16 print(f'PCA reconstruction loss stats based on first 25 PCs, \n (mean,std) =
17     {mean_mse_25, std_mse_25}')
17 print(f'PCA reconstruction loss stats based on first 50 PCs, \n (mean,std) =
18     {mean_mse_50, std_mse_50}')

```

PCA reconstruction loss stats based on first 10 PCs,
 $(\text{mean}, \text{std}) = (523.2417453440711, 257.6412003257671)$

PCA reconstruction loss stats based on first 25 PCs,
 $(\text{mean}, \text{std}) = (332.2564923164667, 153.11014543812558)$

PCA reconstruction loss stats based on first 50 PCs,
 $(\text{mean}, \text{std}) = (198.4252027189417, 84.17954418126885)$

As expected, with the increasing number of PCs used in reconstruction, the error measured between original data space and reconstructions are decreased gradually.

1.3. **Part C.** In this part the question, instead of PCA, we'll find explanatory variables to capture stimulus properties using independent component analysis (ICA). We'll perform the same analysis with ICA instead of PCA.

Let's discuss the internal structure of ICA algorithm and their use cases. ICA is computational method for separating a multivariate signal into additive subcomponents. This is done by assuming that the subcomponents are non-Gaussian signals and that they are statistically independent from each other [2]. These settings are necessary and sufficient for encoding source signals and their mixing properties for ICA. ICA seeks for orthogonal rotation of pre-whitened data with iterative schedule that maximizes the non-Gaussianity of the rotated components [2]. Here, non-Gaussianity property is required for finding orthogonal directions and serves as a proxy for independence in the statistical sense. Whitening is required for ICA computations. Whitening is a statistical transformation in such a way that potential correlations between its components are removed (covariance equal to 0) and the variance of each component is equal to 1. Once the data is whitened, we can perform ICA analysis.

Let's dive into analytical formulation of ICA. In the Sparse Coding settings, one wanted to learn a latent representation as in the case of neurobiological settings. In ICA, we want to learn linearly represented subspace with strict orthonormality properties. More precisely, given the data X , ICA seeks for a set of basis vectors that is represented in the matrix W such that features are sparse and the representation is orthonormal. Hence, with these analytical properties, we can construct objective function $J(W)$ as follows.

$$(6) \quad J(W) = \|WX\|_1 \text{ s.t. } WW^T = I$$

Hence, the ICA problem turns the optimization problem as in the all cases of ML pipelines. There are various ways to optimize the objective function but generally there is no simple analytical solution as usually in the case in ML. Moreover, orthonormality property also make the optimization problem harder. For example, $J(W)$ can be optimized with gradient descent based algorithms but the gradient in the direction of convergence must be followed by a step that maps the new basis back to the space of orthonormal bases. It is feasible but generally slow. Also, there are non-iterative methods for computing ICA that are slightly more complicated than iterative settings. However, the main idea is same. Anyway, these are not the concepts of this homework, let's back to our problem.

There are bunch of ICA algorithms with different degree of computational complexities. In these work, we'll work with the FastICA algorithm to find explanatory variables of face images. Here is the code for computation of FastICA with scikit-learn library.

```
1 fastIca_10 = FastICA(n_components = 10, random_state = 5)
2 fastIca_components_10 = fastIca_10.fit_transform(faces)
```

With these two lines of Python code, we whitened the data and perform FastICA analysis on face images with the dimensionality 10. Let's visualize the ICs as follows.

```
1 fig, axes = plt.subplots(2, 5, figsize=(10,4))
2 for i, ax in enumerate(axes.flat):
3     ax.imshow(fastIca_10.components_[i].reshape(32, 32).T, cmap = 'gray')
4     ax.axis('off')
```



FIGURE 9. Visualization of 10 ICs

FastICA algorithm capture the spatial semantics of face images with 10 components. Let's see further visualization of components with higher degree of ICs as follows.

```

1 fastIca_25 = FastICA(n_components = 25,whiten = True, random_state = 5)
2 fastIca_components_25 = fastIca_25.fit_transform(faces)
3
4
5 fig, axes = plt.subplots(5, 5, figsize=(10,10))
6 for i, ax in enumerate(axes.flat):
7     ax.imshow(fastIca_25.components_[i].reshape(32, 32).T, cmap = 'gray')
8     ax.axis('off')
```



FIGURE 10. Visualization of 25 ICs

As we can see, with the increasing number of IC components, the spatial structure of face images are much more clear and intuitive. Let's see more with higher number of ICs.

```

1 fastIca_50 = FastICA(n_components = 50, random_state = 5)
2 fastIca_components_50 = fastIca_50.fit_transform(faces)
3
4
5 fig, axes = plt.subplots(5, 10, figsize=(20,10))
6 for i, ax in enumerate(axes.flat):
7     ax.imshow(fastIca_50.components_[i].reshape(32, 32).T, cmap = 'gray')
8     ax.axis('off')

```



FIGURE 11. Visualization of 50 ICs

Then, the next part is ICA reconstructions based on the 10,25 and 50 IC components.

```

1 faces_ICA_10 = fastIca_10.inverse_transform(fastIca_components_10)
2 faces_ICA_25 = fastIca_25.inverse_transform(fastIca_components_25)
3 faces_ICA_50 = fastIca_50.inverse_transform(fastIca_components_50)

```

Here, I back to original data space with low rank ICA reconstruction based on 10, 25 and 50 ICs. From now on, I visualize the reconstructions sequentially starting from the next page.

```
1 plot_faces(faces_ICA_10, suptitle = 'FastICA reconstruction based on 10  
→ independent components')
```

FastICA reconstruction based on 10 independant components

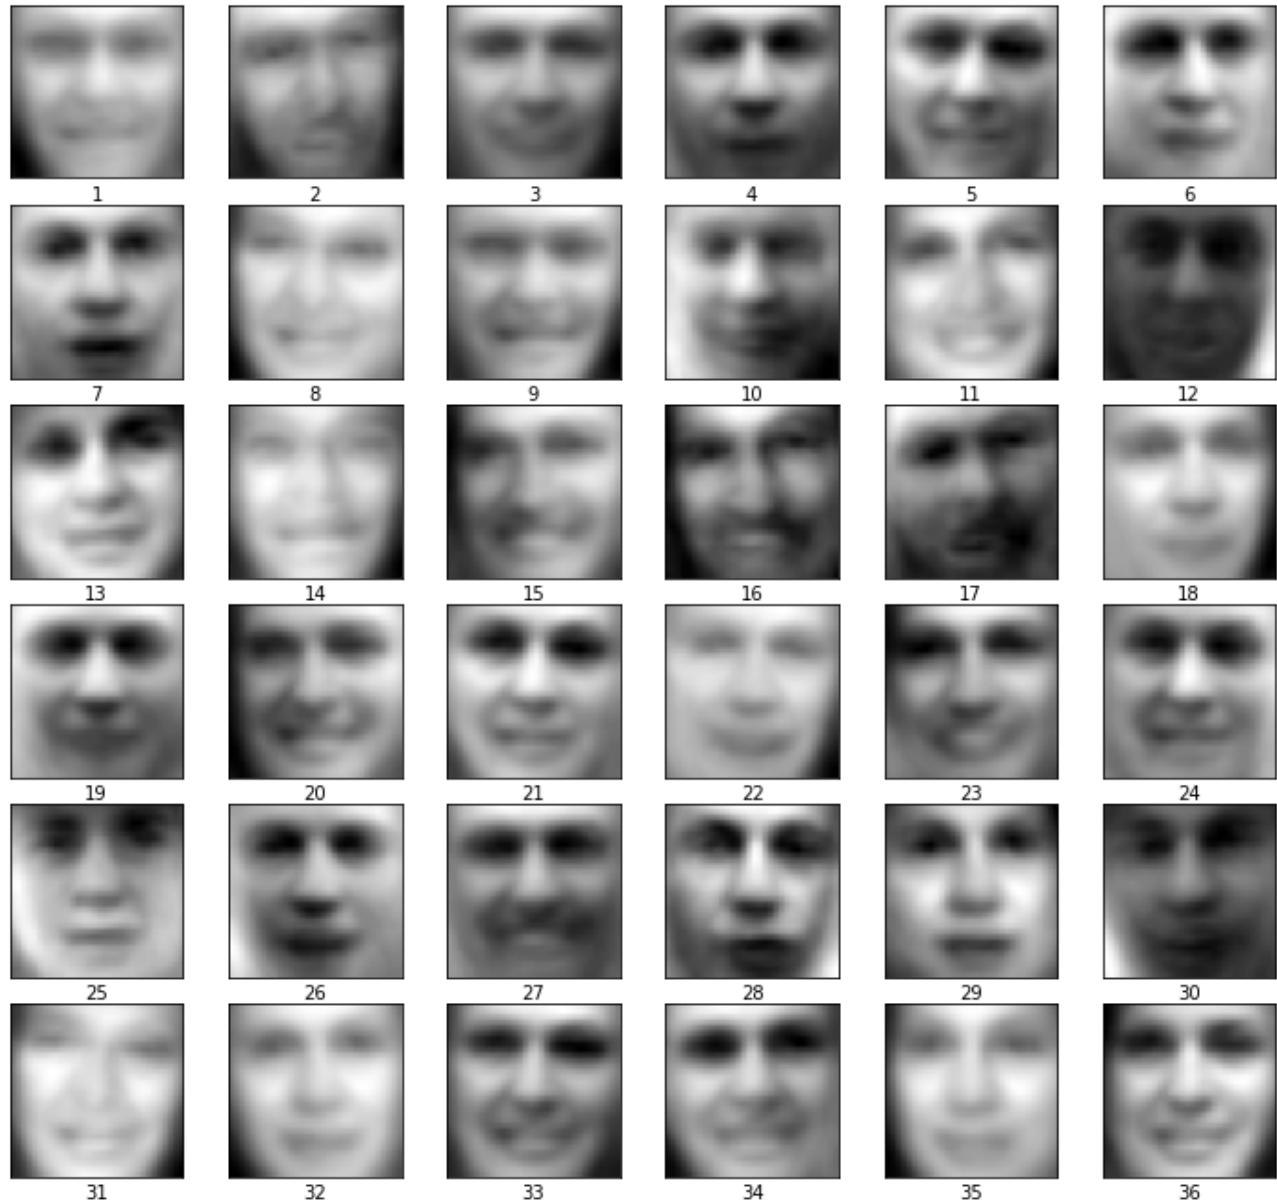


FIGURE 12. FastICA reconstruction based on 10 independent components

As we can see, we capture the semantically meaningful spatial structure of the face images based on the first 10ICs. The spatial densities are clear to identify by human vision. Further, generalized micro mimics of faces are also come to the fore.

```
1 plot_faces(faces_ICA_25, suptitle = 'FastICA reconstruction based on 25  
→ independent components')
```

FastICA reconstruction based on 25 independant components



FIGURE 13. FastICA reconstruction based on 25 independent components

As we can see, we further capture the semantically meaningful spatial structure of the face images based on the first 25 ICs. The spatial densities are much more clear than before. As expected, with increasing dimension of ICA space, we further reason the facial cognitives.

```
1 plot_faces(faces_ICA_50, suptitle = 'FastICA reconstruction based on 50  
→ independent components')
```

FastICA reconstruction based on 50 independant components



FIGURE 14. FastICA reconstruction based on 50 independent components

As a last reconstruction, based on 50 ICs, the algorithm captures the greatest semantic in original feature dimensionality as it can be seen from the spatial similarity between face images and its reconstructions.

Then, let's reason the reconstructions mathematically based MSE loss and its descriptive statistics.

```

1 mse_10 = squared_error(y_true = faces, y_pred = faces_ICA_10)
2 mse_25 = squared_error(y_true = faces, y_pred = faces_ICA_25)
3 mse_50 = squared_error(y_true = faces, y_pred = faces_ICA_50)
4
5
6 std_mse_10 = mse_10.mean(-1).std()
7 std_mse_25 = mse_25.mean(-1).std()
8 std_mse_50 = mse_50.mean(-1).std()
9
10 mean_mse_10 = mse_10.mean()
11 mean_mse_25 = mse_25.mean()
12 mean_mse_50 = mse_50.mean()
13
14
15 print(f'ICA reconstruction loss stats based on first 10 ICs, \n (mean,std) =
    ↪ {mean_mse_10, std_mse_10}')
16 print(f'ICA reconstruction loss stats based on first 25 ICs, \n (mean,std) =
    ↪ {mean_mse_25, std_mse_25}')
17 print(f'ICA reconstruction loss stats based on first 50 ICs, \n (mean,std) =
    ↪ {mean_mse_50, std_mse_50}')

```

ICA reconstruction loss stats based on first 10 ICs,
 $(\text{mean}, \text{std}) = (523.2417453440557, 257.6412004632383)$
ICA reconstruction loss stats based on first 25 ICs,
 $(\text{mean}, \text{std}) = (332.2564920665104, 153.11028826729745)$
ICA reconstruction loss stats based on first 50 ICs,
 $(\text{mean}, \text{std}) = (198.42506719787565, 84.17996584442903)$

First things first, FastICA algorithm is internally structured as whitening and orthogonal rotation as rotation yields the greatest non-Gaussianity with rotated direction. Since the whitening process can be accomplished by classical PCA and orthogonal rotation does not affect the reconstruction MSE loss of the ICA, the MSE loss statistics are identical.

1.4. Part D. In this section, we'll perform explanatory face image analysis with Non-Negative Matrix Factorization (NNMF) instead of PCA or ICA. The computational procedures are identical except for the dimensionality reduction algorithm as the previous sections. In other words, we'll perform NNMF analysis based on 10, 25 and 50 MFs. We'll perform comprehensive visualizations of MFs with their reconstructions to the original data space. Finally, we'll report error statistics with comparisons to other sections. As usual, let's start with the discussion of the algorithm and its use cases.

NNMF is an unsupervised machine learning algorithm and widely used tool to analyze high dimensional datasets as in the case of PCA and ICA except for non-negativity condition as prior. It extract sparse and meaningful insights from high dimensional non-negative features spaces. It approximates a non-negative low rank data matrix such that $X \approx WH$. Hence, given the data matrix $X_{m \times n}$ where the entries of A is strictly non-negative, NNMF learns the $W_{m \times k}$ and $H_{k \times n}$ such that it gives the estimation of A , i.e., $X \approx WH$ where the quantity k is set by the researcher as it denotes the latent dimensionality of the original feature space. Note that W and H are generally called dictionary (or basis) and expansion (or coefficient) matrix, respectively.

Then, let's formulate the NNMF problem as iterative optimization problem as follows.

$$(7) \quad \min_{W \in \mathbb{R}^{m \times k}, H \in \mathbb{R}^{k \times n}} \|X - WH\|_F^2 \text{ s.t. } W_{i,j} \geq 0 \text{ and } H_{j,k} \geq 0 \forall i, j, k$$

where F denotes Frobenius norm that gives the intuition that there are Gaussian noise. Note that another norms can be used depending on the application. (Kullback-Leibler divergence for text-mining, the Itakura-Saito distance for music analysis, or the l_1 norm to improve robustness against outliers).

Furthermore, NNMF learning is a subclass of NP-hard theoretically, but there are heuristic approximations methods works well in practise. So, the the parameters of NNMF, W and H are not unique. Here, one well-known update rule for W and H as it converges to local minima in the error surface is presented.

$$(8) \quad H_{\alpha u} \leftarrow H_{\alpha u} \sum_i W_{i\alpha} \frac{X_{iu}}{(WH)_{iu}} \text{ and } W_{i\alpha} \leftarrow W_{i\alpha} \sum_u H_{\alpha u} \frac{X_{iu}}{(WH)_{iu}}$$

The iteration number is also set by the user before hand. As a final note, as PCA is not localized, oriented to spatial structure of images, NNMF learns a parts-based representation of faces whereas PCA learn holistic representations of images.

Let's move into our case. I set the iteration number as 500 as I observe that this number is a sufficient trade off between accuracy of estimation and its computational complexity. Let's see the code for computation of NNMF and its first 10 components visualization. Note that I also add the absolute value of the minimum of the faces images to shift pixel intensities to positive spatial axis.

```

1 max_iter = 500
2 faces += np.abs(np.min(faces))
3 NMF_10 = NMF(n_components = 10, max_iter = max_iter)
4 NMF_components_10 = NMF_10.fit_transform(faces)
5
6
7 fig, axes = plt.subplots(2, 5, figsize=(10,4))
8 for i, ax in enumerate(axes.flat):
9     ax.imshow(NMF_10.components_[i].reshape(32, 32).T, cmap = 'gray')
10    ax.axis('off')
```



FIGURE 15. Visualization of NNMF 10 MFs

Let's also visualize the further MFs as follows.

```
1 NMF_25 = NMF(n_components = 25, max_iter = max_iter)
2 NMF_components_25 = NMF_25.fit_transform(faces)
3
4
5 fig, axes = plt.subplots(5, 5, figsize=(10,10))
6 for i, ax in enumerate(axes.flat):
7     ax.imshow(NMF_25.components_[i].reshape(32, 32).T, cmap = 'gray')
8     ax.axis('off')
```



FIGURE 16. Visualization of NNMF 25 MFs

```

1 NMF_50 = NMF(n_components = 50, max_iter = max_iter)
2 NMF_components_50 = NMF_50.fit_transform(faces)
3
4
5 fig, axes = plt.subplots(5, 10, figsize=(12,6))
6 for i, ax in enumerate(axes.flat):
7     ax.imshow(NMF_50.components_[i].reshape(32, 32).T, cmap = 'gray')
8     ax.axis('off')

```

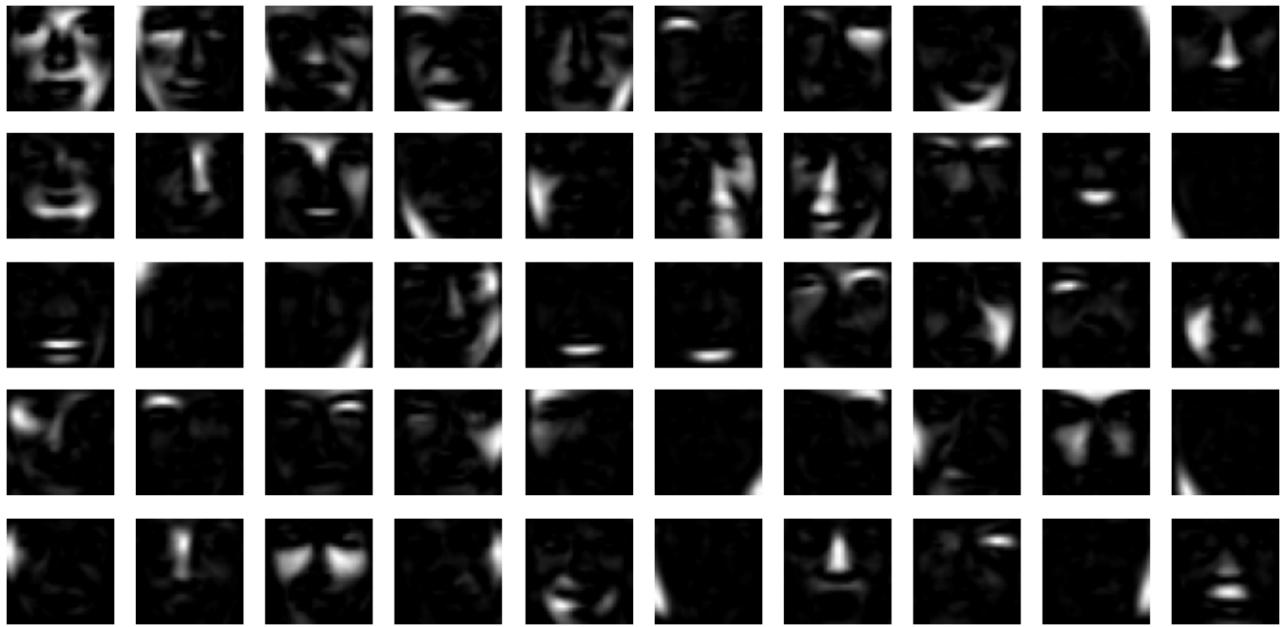


FIGURE 17. Visualization of NNMF 50 MFs

As we can see, with the increasing number of MFs, we capture different spatial information regarding the face images. Some of the components focused on the orientations, while others focused on facial keypoints, face structures etc.

Let's move on the reconstructions based on 10, 25, and 50 MFs with their corresponding MSE loss statistics as follows.

```

1 faces = h5py.File('hw4_data1.mat','r')['faces'][:].T
2 faces_NNMF_10 = NMF_10.inverse_transform(NMF_components_10) -
    ↪ np.abs(np.min(faces))
3 faces_NNMF_25 = NMF_25.inverse_transform(NMF_components_25) -
    ↪ np.abs(np.min(faces))
4 faces_NNMF_50 = NMF_50.inverse_transform(NMF_components_50) -
    ↪ np.abs(np.min(faces))

```

Note that the absolute value of the minimum of the faces are subtracted from the reconstructions since we added these terms as we are preparing faces images to the NNMF algorithm. Also note that, since I did not use another variable for added version of face images, I read the face image data again, then subtract the the absolute value of the minimum values of the face image data data.

```
1 plot_faces(faces_NNMF_10, suptitle = 'NNMF Reconstruction of faces based on  
→ 10MFs')
```

NNMF Reconstruction of faces based on 10MFs



FIGURE 18. NNMF Reconstruction of faces based on 10 MFs

```
1 plot_faces(faces_NNMF_25, suptitle = 'NNMF Reconstruction of faces based on  
→ 25MFs')
```

NNMF Reconstruction of faces based on 25MFs



FIGURE 19. NNMF Reconstruction of faces based on 25 MFs

```
1 plot_faces(faces_NNMF_50, suptitle = 'NNMF Reconstruction of faces based on  
→ 50MFs')
```

NNMF Reconstruction of faces based on 50MFs



FIGURE 20. NNMF Reconstruction of faces based on 50MFs

Then, let's evaluate our NNMF model based on MSE loss statistics as follows.

```

1 mse_10 = squared_error(y_true = faces, y_pred = faces_NNMF_10)
2 mse_25 = squared_error(y_true = faces, y_pred = faces_NNMF_25)
3 mse_50 = squared_error(y_true = faces, y_pred = faces_NNMF_50)
4
5
6 std_mse_10 = mse_10.mean(-1).std()
7 std_mse_25 = mse_25.mean(-1).std()
8 std_mse_50 = mse_50.mean(-1).std()
9
10 mean_mse_10 = mse_10.mean()
11 mean_mse_25 = mse_25.mean()
12 mean_mse_50 = mse_50.mean()
13
14
15 print(f'NNMF reconstruction loss stats based on first 10 MFs, \n (mean,std) =
    ↪ {mean_mse_10, std_mse_10}')
16 print(f'NNMF reconstruction loss stats based on first 25 MFs, \n (mean,std) =
    ↪ {mean_mse_25, std_mse_25}')
17 print(f'NNMF reconstruction loss stats based on first 50 MFs, \n (mean,std) =
    ↪ {mean_mse_50, std_mse_50}')

```

NNMF reconstruction loss stats based on first 10 MFs,
 (mean,std) = (533.7720128448966, 267.388838812122)
 NNMF reconstruction loss stats based on first 25 MFs,
 (mean,std) = (351.018592518666, 169.342734092105)
 NNMF reconstruction loss stats based on first 50 MFs,
 (mean,std) = (221.34231024320826, 103.62372565762264)

We observe that NNMF reconstruction gives the least performance among other methods. The reconstruction statistics are not much different from the ICA and PCA case, but they were slightly better than NNMF. Let's interpret the results and make a discussion.

PCA model produces a new data features as result of combination of existing one while NNMF just decompose a dataset matrix into its non-negative sub matrix whose dimensionality is uneven. Hence, the output of NNMF can be visualized as compressed version of original dataset while reconstructions of PCA uncover the embedding gradients of the spatial data as we previously see. Hence, NNMF outputs are interpretable. In the settings of PCA, we assume that the data lies in lower linear subspace of the original data space. So, non-linear modifications are necessary for the data by nature while NNMF don't require any modifications irrespective of linear separability. Finally, NNMF is sensitive to initialization, hence produces non-unique outputs in every fit while PCA gives the same PCs with different number of trials. According to the error statistics found based on different unsupervised dimension reduction algorithm, PCA and ICA gives the least MSE error hence performed better among others.

2. QUESTION 2

In this question, we consider a population of 21 independent neurons with Gaussian-shaped tuning curves:

$$(9) \quad f_i(x) = A * e^{-\frac{(x-\mu_i)^2}{2\sigma^2}}$$

The tuning curves have an amplitude of 1 and a standard deviation of $\sigma_i = 1$, with centers μ_i evenly spaced between -10 and 10 along the x-axis.

2.1. Part A. In this part, we'll plot all tuning curves in the population on the same axis. Then, we'll simulate the population response to the stimulus $x = -1$, and plot the population response as a function each neuron's preferred stimulus value.

Let's dive into action. Here I created a Gauss tuning function that produce Gaussian shaped tuning function of a population of neurons as follows.

```

1 def gauss_tuning(x:np.ndarray = np.linspace(-15, 16, 500),
2                   mu:float = 1,
3                   sigma:float = 1,
4                   A:float = 1) -> np.float16:
5
6     """
7         Gaussian shaped tuning function of a population of neurons.
8
9     Arguments:
10        x      (np.ndarray)    : The input stimulus parameters
11        A      (float)        : Gain of the Gaussian-shaped tuning curve
12        mu     (float)        : Mean of the Gaussian-shaped tuning curve
13        sigma  (float)        : Standard deviation of the Gaussian-shaped
14        ↪      tuning curve
15
16        Returns:
17            response : Resulting neural response
18        """
19
20    return A * np.exp(-0.5 * ((x- mu)/sigma) ** 2)

```

Then, I plot the responses as well as their averaged version (as additional) according to hyperparameters given in the question.

```

1 neural_responses = []
2 legends = []
3 stimuli = np.linspace(-15, 16, 500)
4 means = np.arange(-10, 11)
5 plt.figure(figsize=(10,5))
6
7 for mean in means:
8     response = gauss_tuning(mu = mean)
9     plt.plot(stimuli, response)
10    legends.append(f" Response with mean {mean}")
11
12    # Let's keep neural responses for future use:
13    neural_responses.append(response)
14
15    # Plot the tuning profiles

```

```

16 plt.plot(stimuli, np.mean(neural_responses, axis=0), color = '0')
17 legends.append(f" Average Response")
18 plt.legend(legends, loc="right", bbox_to_anchor=(1.4, 0.5))
19 plt.xlabel('Stimulus')
20 plt.ylabel('Activity')
21 plt.title('Tuning Curves of a Population of Neurons')
22 plt.grid()
23 plt.show()

```

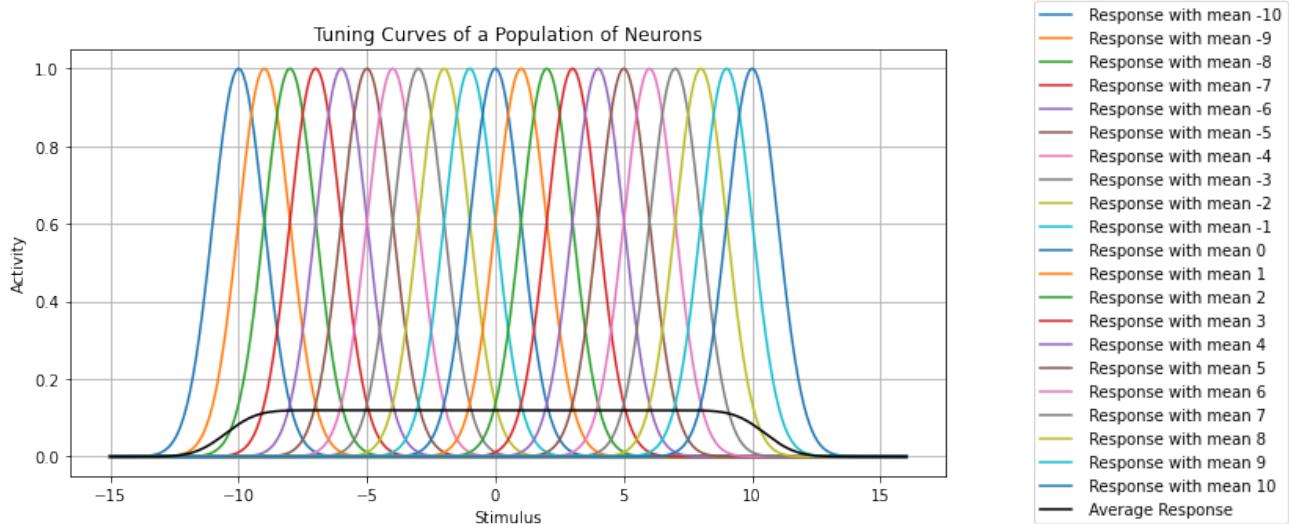


FIGURE 21. Gaussian Tuning Curves of a Population of Neurons

As we can see, originated from the Gaussian, each neurons preferred stimulus is the mean of its Gaussian curve. Then, let's see the simulation of the population response to the stimulus with $x = -1$ as follows.

```

1 kwargs = dict(
2         color = '0',
3         marker= 'o',
4         markerfacecolor='red'
5     )
6 plt.figure(figsize=(10,5))
7 plt.plot(means, gauss_tuning(-1, mu = means),**kwargs)
8 plt.xlabel('Preferred Stimulus')
9 plt.ylabel('Population Response')
10 plt.title('Population Response to the Stimulus x = -1 vs Preferred Stimuli of
11 → Neurons')
12 plt.grid()
13 plt.show()

```

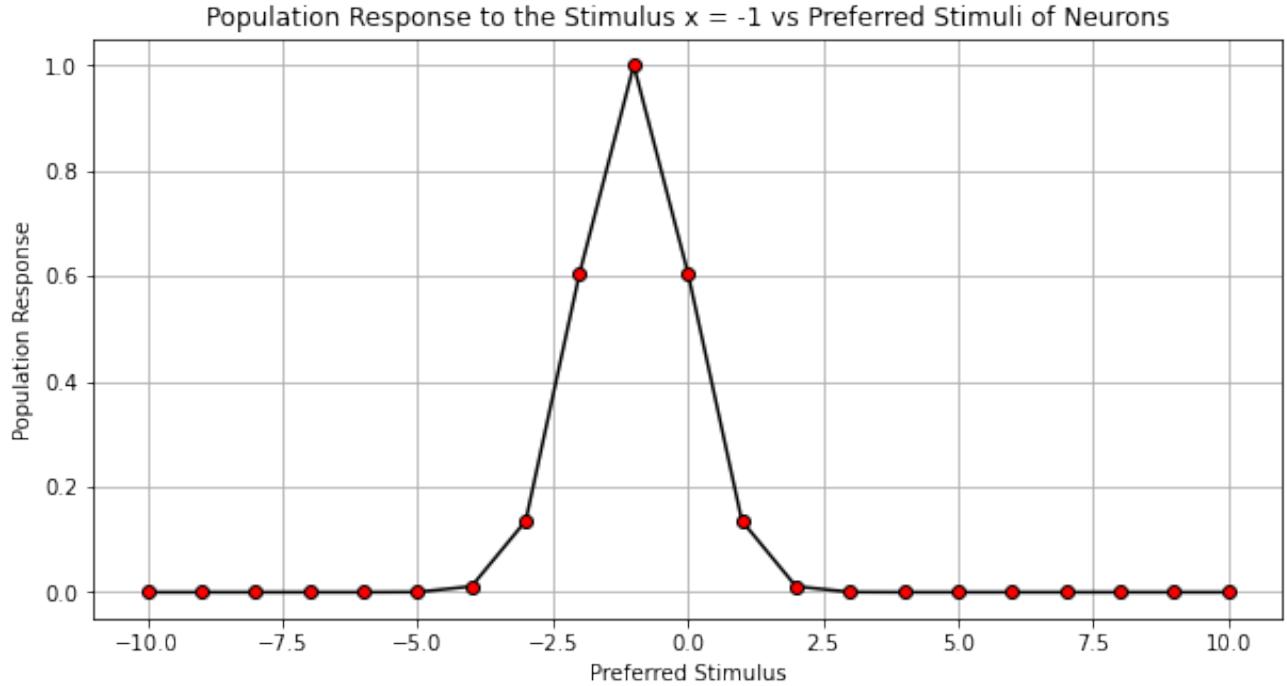


FIGURE 22. Neural Population Response to the Stimulus $x = -1$

As we distributed the neural responses as a Gaussian, preferred stimulus is where the stimulus that elicits a maximal response. It is expected due to the Gaussian properties.

2.2. Part B. In this part of the question, we'll perform a simulated experiment with 200 trials, in each trial we sample a stimulus intensity uniformly from the interval $[-5, 5]$ simulate the 21-long vector of population response. Also, we assume the corruption of neural response described by the centered Gaussian noise with $\sigma/20$. Then, the next is to decode the neural response by the Winter-Take-All decoder. Winner-take-all is a computational principle applied in computational models of neural networks by which neurons in a layer compete with each other for activation [3]. In the classical form, only the neuron with the highest activation stays active while all other neurons shut down; however, other variations allow more than one neuron to be active, for example the soft winner take-all, by which a power function is applied to the neurons [3]. Hence, we'll simply try to find a stimulus that elicits a maximal response among all neural responses. From the computational perspective, the problem is argument maximizer problem as follows.

$$(10) \quad x_{WTA} = \operatorname{argmax}_{u_i} r_i$$

where x_{WTA} is the stimulus that maximizes the response (as mean in the Gaussian case) and r_i is the i^{th} neural response. Let's see in action.

```

1 def WTA_decoder(stimuli:np.ndarray, response:np.ndarray) -> np.float16:
2     """
3         Given a population response and stimuli of the
4         neurons, compute the winner-take-all decoder that
5         estimates the actual stimulus as the preferred
6         stimulus of the neuron with maximum response
7
8     Arguments:
9         stimuli (np.ndarray): The preferred stimuli of the neurons
10        response (np.ndarray): The neural responses
11    Returns:
12        stimulus (np.float16): the estimated input stimulus that maximizes
13        → the response
14        """
15
16        response += np.random.normal(loc = 0, scale = 1/20, size = (21,))
16    return stimuli[np.argmax(response)]

```

As we can see, it is simply argument maximizer operation that finds a stimulus that elicits a maximal neural response. Let's see visualization of the performance of WTA decoder with error statistics as follows.

```

1 n_trials = 200
2 stimuli_interval = np.linspace(-5,5, 500).tolist()
3
4 # To keep 'responses', 'stimuli', 'WTA_stimuli', 'Errors'
5 stimuli_response_ = []
6
7 for stimuli in random.sample(stimuli_interval, n_trials):
8     response = gauss_tuning(stimuli, mu = means)
9     WTA_stimuli = WTA_decoder(means, response)
10    stimuli_response_.append((response, stimuli, WTA_stimuli, np.abs(WTA_stimuli
11        - stimuli)))
11
12 # Tuples are gathered:
13 stimuli_response = list(zip(*stimuli_response_))

```

Then, here is the visualization.

```

1 fig = plt.figure(figsize=(10,5))
2 error = np.array(stimuli_response[-1])
3 fig.suptitle(f'Error Stats: (mean,std) =
4     {round(error.mean(),2),round(error.std(),2)}')
4 plt.scatter(range(n_trials), stimuli_response[1], marker="o", color="red", s=30,
5     linewidths=1)
5 plt.scatter(range(n_trials), stimuli_response[2], marker="x", color="green",
6     s=30, linewidths=1)
6 plt.xlabel('# of Trials')
7 plt.ylabel('Stimulus')
8 plt.title('Actual and WTA Estimated Stimuli Across Trials')
9 plt.legend(['Actual', 'WTA Estimated'], loc='upper right')
10 plt.show()

```

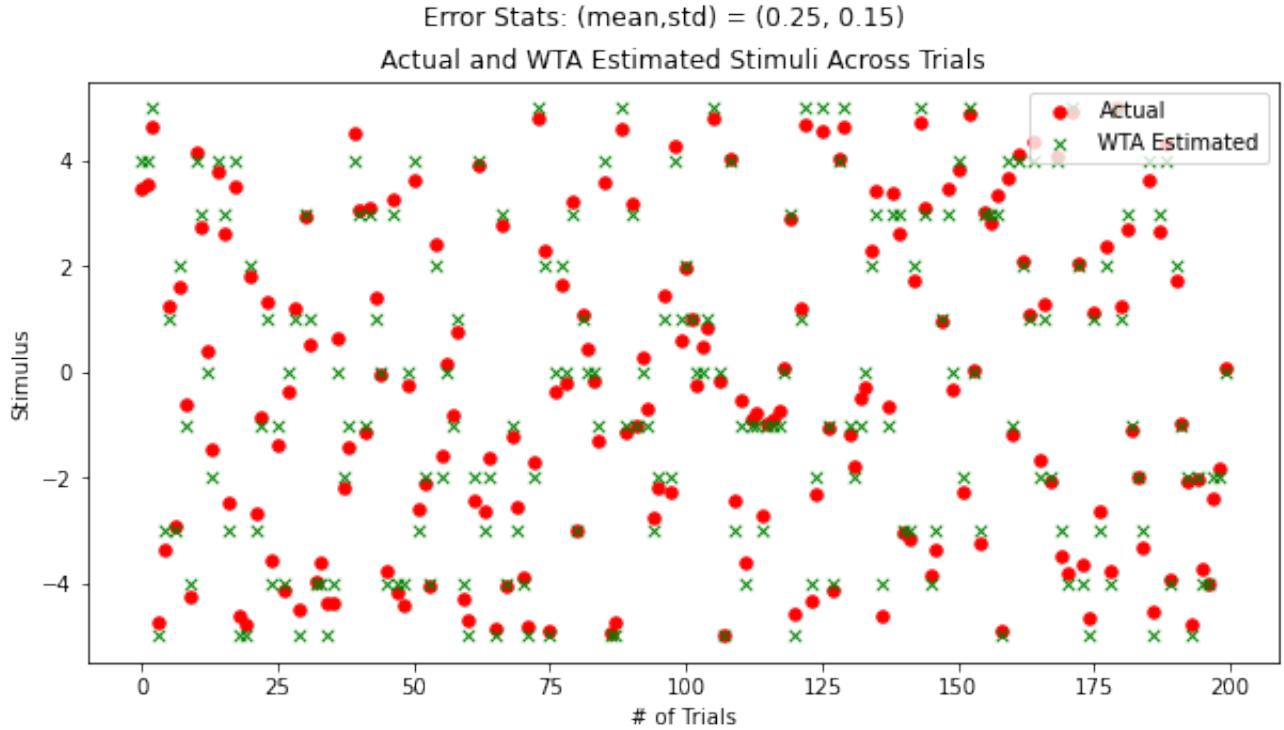


FIGURE 23. Actual and WTA Estimated Stimuli Across Trials

As we can see, WTA stimulus shots are not bad, we are closer to actual stimulus. Let's see the error statistics for future comparison.

```

1 print('Error Statistics for Winner Take All Decoder')
2 print('Mean of errors in stimuli estimation:', error.mean().round(5))
3 print('Standard deviation of errors in stimuli estimation :',
      error.std().round(5))

```

Error Statistics for Winner Take All Decoder
 Mean of errors in stimuli estimation: 0.26058
 Standard deviation of errors in stimuli estimation : 0.15047

Let's move to other decoding algorithms and then compare our results with WTA decoder.

2.3. Part C. In this part, for the same experimental trials simulated in part b, we'll implement a maximum-likelihood decoder, and calculate the stimulus estimate x_{ML} for each trial. As we did in part b, we'll provide estimation error statistics with corresponding visualizations. Let's derive the analytical parts, and apply the trick of logarithmic differentiation to decode neural responses.

Let x_{ML} estimation of the stimulus by MLE algorithm, r_i is the i^{th} response for $i = 1, \dots, 21$ that represents the 21 neurons responses. Then, ML decoder try to find a x_{ML} as follows.

$$(11) \quad x_{ML} = \operatorname{argmax}_x P(r_1, \dots, r_{21} | x)$$

The responses of neural population r_1, \dots, r_{21} is modeled as follows.

$$(12) \quad r_1, \dots, r_{21} = f_{i=1, \dots, 21}(x) + \mathcal{N}(0, (\sigma/20)^2)$$

Hence, by the property of any linear combination of Gaussian is Gaussian, we have

$$(13) \quad r_1, \dots, r_{21} \sim \mathcal{N}(f_{i=1,\dots,21}, (\sigma/20)^2)$$

The rest is just algebra, here is quick proof of end result.

$$P(r_1, \dots, r_{21} | x) = \prod_{i=1}^{21} \log \mathcal{N}(f_{i=1,\dots,21}, (\sigma/20)^2)$$

We simply take logarithm of the equation as MLE trick to simplify analytical derivations of probability distribution differentiations. So, when we talk about the MLE of a sample, a product naturally arises because the joint distribution of independent observations (x_1, \dots, x_n) is given by the product of the marginal distributions of each observation; i.e.,

$$(14) \quad P(r_1, \dots, r_{21} | x) = \prod P(r_i | x).$$

So to find the maximum likelihood, it is usually easier to apply a monotone transformation to the likelihood (thus preserving the location of relative extrema) that converts multiplication to addition.

Then, with little bit of mathematical manipulation, $P(r_1, \dots, r_{21} | x)$ becomes

$$(15) \quad P(r_1, \dots, r_{21} | x) \underset{\approx}{\sim} - \sum_{i=1}^{21} (r_i - f_i)^2$$

Then, we just put the derivation into optimization problem format as follows.

$$(16) \quad x_{ML} = \underset{x}{\operatorname{argmax}} P(r_1, \dots, r_{21} | x) = \underset{x}{\operatorname{argmin}} \left(\sum_{i=1}^{21} (r_i - f_i)^2 \right)$$

Note that the symbol x was thrown from the notation of $r_i(x)$ (become r_i) and $f_i(x)$ (become f_i) for notation simplification. Hence, we derive the analytical formulation of the MLE decoder, let's see the Python code for that in the following page.

```

1 def MLE_decoder(stimuli_interval:np.ndarray = np.linspace(-5,5, 500).tolist(),
2                 response:np.ndarray = None) -> np.float16:
3
4     """
5         Given a population response and stimuli of the
6         neurons, compute the MLE decoder that
7         estimates the actual stimulus as the preferred
8         stimulus of the neuron with maximum response
9
10    Arguments:
11        stimuli (np.ndarray): The preferred stimuli of the neurons
12        response (np.ndarray): The neural responses
13
14    Returns:
15        stimulus (np.float16): the estimated input stimulus that maximizes
16        the response
17    """
18    logs = list()
19
20    for stim in stimuli_interval:
21        log = sum((response_ - gauss_tuning(stim, mu_)) ** 2 for response_, mu_
22                   in zip(response, means))
23        logs.append(log)
24    idx_stim_max = np.argmin(logs)
25    return stimuli_interval[idx_stim_max]

```

Hence, we just plug the analytical form to our function that decode the neural response. Hence, we can move to simulation and error statistics with corresponding visualization as follows.

```

1 MLE_logs_ = list()
2 for neural_activity in stimuli_response_:
3     response, stimulus = neural_activity[:2]
4     estimated_stimulus_MLE = MLE_decoder(stimuli_interval, response)
5     MLE_logs_.append((stimulus, estimated_stimulus_MLE, np.abs(stimulus -
6                           estimated_stimulus_MLE)))
7
8 # Tuples are gathered:
9 MLE_logs = list(zip(*MLE_logs_))
10
11 fig = plt.figure(figsize=(10,5))
12 error = np.array(MLE_logs[2])
13 fig.suptitle(f'Error Stats: (mean,std) =
14   {round(error.mean(),2),round(error.std(),2)}')
15 plt.scatter(range(n_trials), MLE_logs[0], marker="o", color="r", s=30,
16             linewidths=1)
17 plt.scatter(range(n_trials), MLE_logs[1], marker="x", color="green", s=30,
18             linewidths=1)
19 plt.xlabel('# of Trials')
20 plt.ylabel('Stimulus')
21 plt.title('Actual and MLE Estimated Stimuli Across Trials')
22 plt.legend(['Actual', 'MLE Estimated'], loc='upper right')
23 plt.show()

```

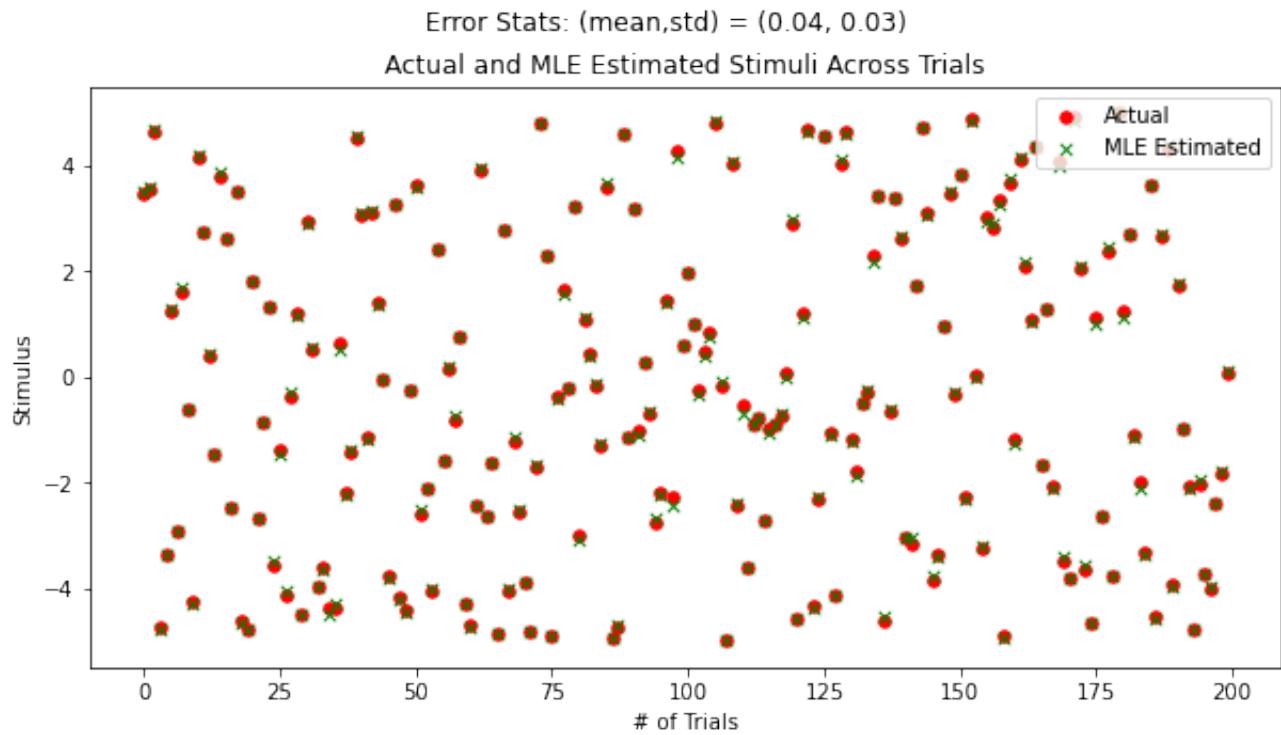


FIGURE 24. Actual and MLE Estimated Stimuli Across Trials

Intuitively, from the figure, MLE outperforms the WTA decoder as its shots are more accurate. Let's evaluate the model from its mean absolute error based statistics as follows.

```

1 print('Error Statistics for MLE Decoder')
2 print('Mean of errors in stimuli estimation:', error.mean().round(5))
3 print('Standard deviation of errors in stimuli estimation :',
      error.std().round(5))

```

Error Statistics for MLE Decoder

Mean of errors in stimuli estimation: 0.04098

Standard deviation of errors in stimuli estimation : 0.03074

According to the comparison of error statistics between WTA and MLE decoder, MLE decoder is clear winner as it decreases our error rates more than %5 in average.

2.4. Part D. In this part, we'll make same experiment and simulation as in the part c except for its decoding mechanism. We replace MLE decoder to MAP decoder. MLE and MAP counterparts in the context of probability theory and statistics. MLE treats the unknown variable as unknown constant, and it falls into the classical statistics concept whereas MAP considers unknown variable as random variable with prior distribution and it falls into Bayesian statistics. Hence, in the Bayesian view, they are treated as random variables with known prior. In classical view, they are treated as deterministic quantities that happen to be unknown.

In the settings of the question, we assume that the prior of the stimulus value x follows a Gaussian distribution with a mean of 0 and a standard deviation of 2.5.

$$P(x | r_1, \dots, r_{21}) \underset{\sim}{\propto} P(r_1, \dots, r_{21} | x)P(x)$$

$$x_{MAP} = \underset{x}{\operatorname{argmax}} P(r_1, \dots, r_{21} | x)P(x) \text{ where } P(x) \sim \mathcal{N}(\mu = 0, \sigma = 2.5)$$

We kindly drop the normalization constant in the posterior calculation of $P(x | r_1, \dots, r_{21})$ since it does not give any further information regarding the estimation of the stimulus except for normalization.

Hence, we can expect that our MAP decoder gives similar analytical expression, but not exactly same. With the MAP decoder settings, we have priory Gaussian denoted by $P(x)$ that affects the analytical derivation. It just adds an extra penalty/regularization term to the analytical expression to the decoder. Since the calculations are very similar to MLE settings except for priory Gaussian aggregation, we directly go to the implementation part as follows.

```

1 def MAP_decoder(stimuli_interval:np.ndarray = np.linspace(-5,5, 500).tolist(),
2                  response:np.ndarray = None) -> np.float16:
3     """
4         Given a population response and stimuli of the
5         neurons, compute the MAP decoder that
6         estimates the actual stimulus as the preferred
7         stimulus of the neuron with maximum response
8
9     Arguments:
10        stimuli (np.ndarray): The preferred stimuli of the neurons
11        response (np.ndarray): The neural responses
12    Returns:
13        stimulus (np.float16): the estimated input stimulus that maximizes
14        → the response
15    """
16    logs = list()
17
18    for stim in stimuli_interval:
19
20        log = sum((r - gauss_tuning(stim, m)) ** 2 for r, m in zip(response,
21                         → means))
22        log = log * 200 + (stim ** 2) / 10
23
24        logs.append(log)
25
26    idx_stim_max = np.argmin(logs)
27    return stimuli_interval[idx_stim_max]
```

We can see that as a additional code to MLE, we simply add the penalty term originated from the priory Gaussian. Let's see the performance of MAP decoder, with its error statistics and prediction visualization.

```

1 MAP_logs_ = list()
2 for neural_activity in stimuli_response_:
3     response, stimulus = neural_activity[:2]
4     estimated_stimulus_MAP = MAP_decoder(stimuli_interval, response)
5     MAP_logs_.append((stimulus, estimated_stimulus_MAP, np.abs(stimulus -
6         ↪ estimated_stimulus_MAP)))
7
8 # Tuples are gathered:
9 MAP_logs = list(zip(*MAP_logs_))
10
11 fig = plt.figure(figsize=(10,5))
12 error = np.array(MAP_logs[2])
13 fig.suptitle(f'Error Stats: (mean,std) =
14     ↪ {round(error.mean(),2),round(error.std(),2)}')
15 plt.scatter(range(n_trials), MAP_logs[0], marker="o", color="r", s=30,
16     ↪ linewidths=1)
17 plt.scatter(range(n_trials), MAP_logs[1], marker="x", color="green", s=30,
18     ↪ linewidths=1)
19 plt.xlabel('# of Trials')
20 plt.ylabel('Stimulus')
21 plt.title('Actual and MAP Estimated Stimuli Across Trials')
22 plt.legend(['Actual', 'MAP Estimated'], loc='upper right')
23 plt.show()

```

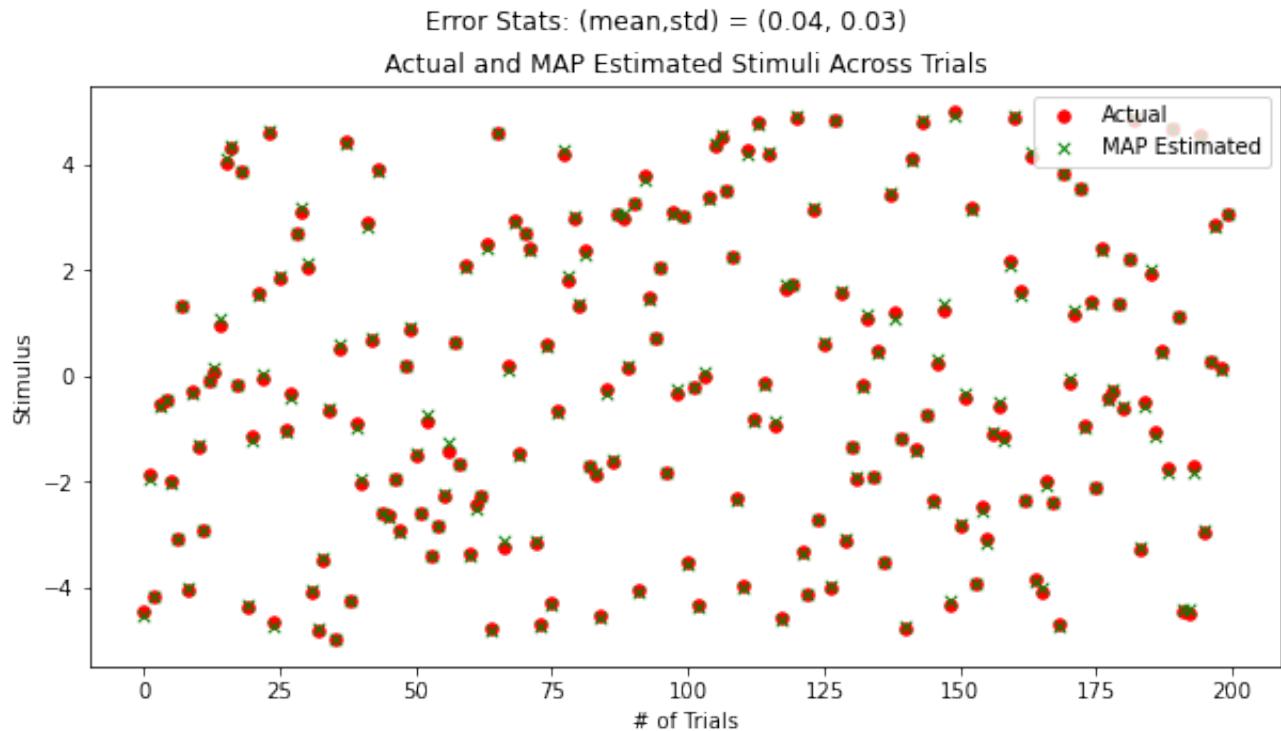


FIGURE 25. Actual and MAP Estimated Stimuli Across Trials

The results are promising. Let's see the error statistics as follows.

```

1 print('Error Statistics for MAP Decoder')
2 print('Mean of errors in stimuli estimation:', error.mean().round(5))
3 print('Standard deviation of errors in stimuli estimation :',
      → error.std().round(5))

```

Error Statistics for MAP Decoder

Mean of errors in stimuli estimation: 0.04068

Standard deviation of errors in stimuli estimation : 0.0311

As we can compare, the results are very close to the MLE case, MAP outperforms the MLE in averaged error case but it fails to win in error standard deviation comparison. Hence, the variance of MAP predictions have less bias than MLE case, but it has higher variance than MLE.

2.5. Part E. In this part, we'll perform an experiment with 200 trials of stimulus intensity. In each trial, sample a stimulus intensity from the interval $[-5, 5]$. For the resulting stimulus vector (of length 200), we'll separately simulate the population response vectors r for $\sigma_i = 0.1$, $\sigma_i = 0.2$, $\sigma_i = 0.5$, $\sigma_i = 1$, $\sigma_i = 2$, and $\sigma_i = 5$. In each case, we assume additive Gaussian noise with zero mean and $1/20$ standard deviation. After that, we'll calculate MLE estimates of the stimulus x_{ML} based on each population response separately.

Since this question is direct computational problem, let's quickly dive into code as follows.

```

1 stds = [0.1, 0.2, 0.5, 1.0, 2.0, 5.0]
2 bias_outer = []
3
4 for stimuli in random.sample(stimuli_interval.tolist(), n_trials):
5     bias_inner = []
6     for std in stds:
7         response = gauss_tuning(stimuli,means,sigma=std)
8         response += np.random.normal(loc = 0, scale = 1/20, size = (21,))
9         stimulus_MLE = MLE_decoder(response=response)
10        bias = np.abs(stimuli - stimulus_MLE)
11        bias_inner.append(bias)
12    bias_outer.append(tuple(bias_inner))
13 # Tuples are gathered:
14 sigmas_response_bias = list(zip(*bias_outer))

```

Then, let's see the error results with table format as follows.

```

1 error_stats = {}
2 for bias, sigma in zip(sigmas_responses,stds):
3     ave_error = np.mean(bias).round(3)
4     std_error = np.std(bias).round(3)
5     error_stats[sigma] = ave_error, std_error
6 error_stats = pd.DataFrame(error_stats).T
7 error_stats.columns = ['Mean Error', 'STD Error']
8 error_stats.index = stds
9 error_stats.index.name = 'Sigma Values'
10 error_stats.sort_values('Mean Error', ascending = True,inplace = True)
11 error_stats.head(n = 6)

```

Here is the table representation of error statistics.

	Mean Error	STD Error
Sigma Values		
1.0	0.044	0.032
0.5	0.055	0.039
2.0	0.092	0.062
5.0	0.371	0.280
0.2	0.515	1.313
0.1	1.877	2.371

TABLE 1. Error Statistics based on σ values with MLE

We can see that we have optimal σ value at $\sigma = 1.0$. Moreover, there is no linear pattern that σ values reasons to error. Let's see the distribution of σ values on mean and standard deviation of errors.

```

1 fig, ax = plt.subplots(figsize = (10,5))
2 error_stats.plot.kde(ax=ax)
3 error_stats.plot.hist(density=True, ax = ax)
4 ax.set_ylabel('Error Stats')
5 ax.set_xlabel('Sigma values $\sigma$')
6 ax.grid()
7
8 fig, ax = plt.subplots(figsize = (10,5))
9 error_stats.plot(ax=ax,marker = 'o')
10 ax.set_ylabel('Error Stats')
11 ax.set_xlabel('Sigma values $\sigma$')
12 ax.grid()
```

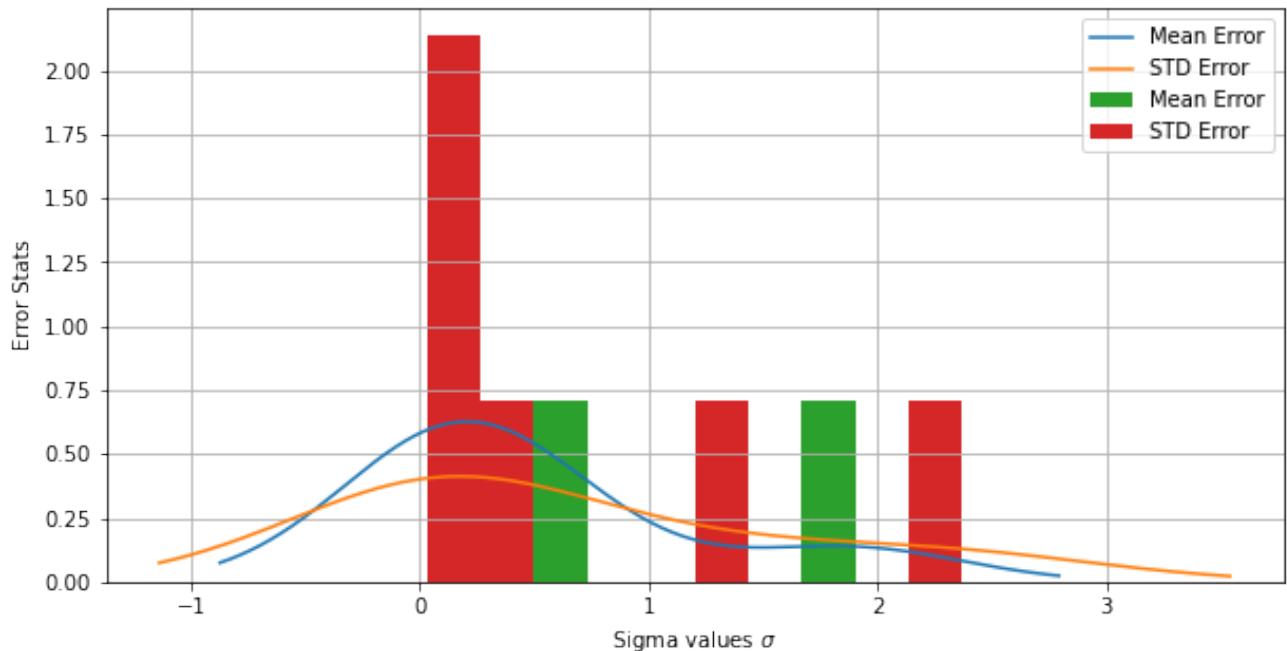


FIGURE 26. Histogram and Density Plot of Error Rate w.r.t. varying σ

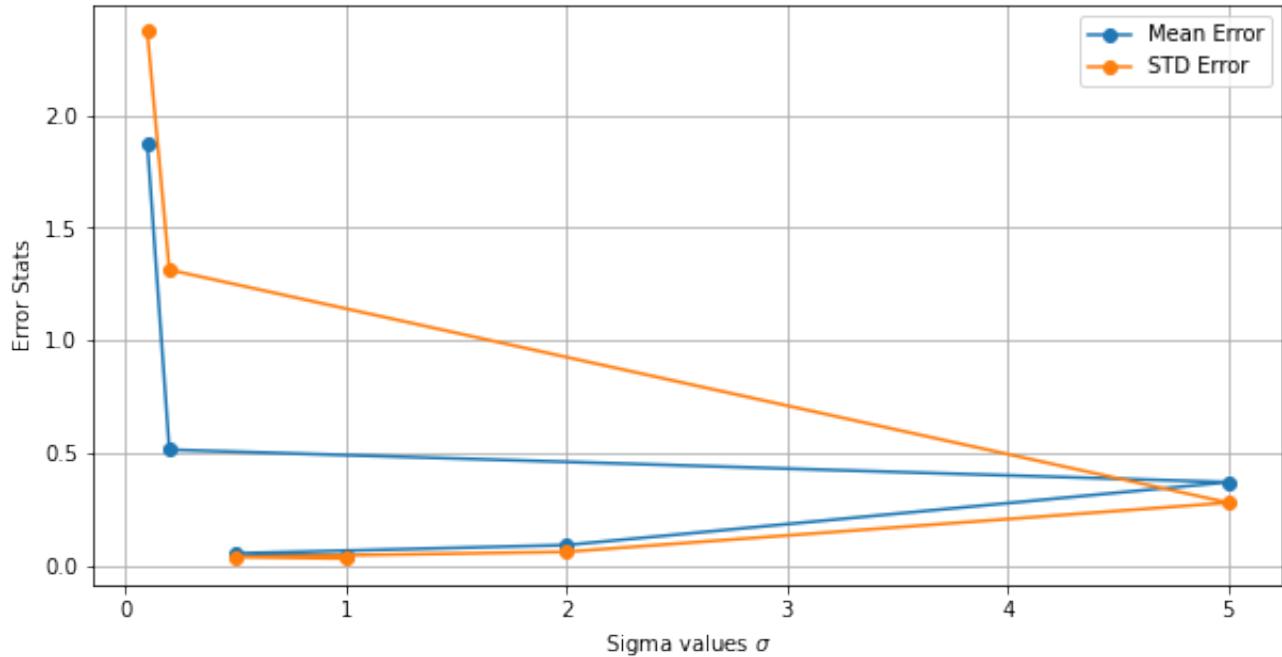


FIGURE 27. Plot of Error Rate w.r.t. varying σ

From the figures, we can conclude that there is no linear pattern of error rates w.r.t. varying parameter σ . But, we can see that we have optimal σ value at $\sigma = 1.0$ that minimizes the error rates based on MLE decoder. Since there is no direct correlation between the narrowness of the Gaussian and error rates, one cannot generalize the effects of σ of the Gaussian tuning with error rates based on the MLE decoder.

3. SOURCE CODE

```
1 #!/usr/bin/env python
2 # coding: utf-8
3 # In[4]:
4
5
6 # Imports:
7 import numpy as np, matplotlib.pyplot as plt, scipy.stats as stats, pandas as pd
8 from sklearn.decomposition import PCA, FastICA, NMF
9 import random, h5py
10
11 settings = np.seterr(all='ignore')
12
13
14 # #### Part A
15
16 # In[5]:
17
18
19 # Retrieving data:
20 faces = h5py.File('hw4_data1.mat','r')['faces'][:,].T
21
22 print(faces.shape)
23
24
25 # In[6]:
26
27
28 # Little bit of dimension manipulation for representing images:
29 N, num_pixel = faces.shape
30 image_faces = faces.reshape(N, np.int(np.sqrt(num_pixel)),
31                             ↳ np.int(np.sqrt(num_pixel)))
31 print(image_faces.shape)
32
33
34 # In[7]:
35
36
37 # Let's look at the face images:
38 fig, axs = plt.subplots(3,3,figsize = (8,8))
39 for i, axes in enumerate(axs.flatten()):
40     axes.imshow(image_faces[i], cmap = 'gray')
41     axes.axis('off')
42
43
44 # In[8]:
45
46
47 # Latent representation dimension:
48 latent_dim = 100
49 pca = PCA(n_components = latent_dim)
50 principalComponents = pca.fit_transform(faces)
```

```

51
52 num2str = lambda x : str(round(sum(x),3))
53
54 legends = [
55
56     pca.explained_variance_ratio_[:10],
57     pca.explained_variance_ratio_[:25],
58     pca.explained_variance_ratio_[:50],
59     pca.explained_variance_ratio_[:]
60
61 ]
62
63 legends = list(map(num2str,legends))
64
65 plt.figure(figsize = (10,5))
66 plt.plot(pca.explained_variance_ratio_, color = 'r')
67 plt.xlabel('PCs')
68 plt.ylabel('Explained Variance')
69 title = 'Principal Components versus Explained Variance \n'
70 title += 'First (10,25,50,100) PCs explained variance = '
71 title += '(' + ' '.join(legends) + ')'
72 plt.title(title)
73 plt.grid()
74 plt.show()
75
76
77 pca_logs = f'Variance explained by PCA for 10 components {legends[0]}\n'
78 pca_logs += f'Variance explained by PCA for 25 components {legends[1]}\n'
79 pca_logs += f'Variance explained by PCA for 50 components {legends[2]}\n'
80 pca_logs += f'Variance explained by PCA for 100 components {legends[3]}'
81 print(pca_logs)
82
83
84 # In[9]:
85
86
87 fig, axes = plt.subplots(5, 5, figsize=(12,12))
88 for i, ax in enumerate(axes.flat):
89     ax.imshow(pca.components_[i].reshape(32, 32).T, cmap = 'gray')
90     ax.axis('off')
91
92
93 # ### Part B
94
95 # In[10]:
96
97
98 def pca_reconstruction(data:np.ndarray,
99                         trained_pca:PCA,
100                        number_PCs:int) -> np.ndarray:
101     """
102
103         Given the input data, trained PCA variable and # of PCs, reconstruct
104         → images based on the PCs components.

```

```

105
106     Arguments:
107         - data (np.ndarray) : Input data
108         - trained_pca (PCA) : trained PCA variable
109         - number_PCs (int) : # of PCs to reconstruct images
110
111     Returns:
112         - reconstructed_data (np.ndarray) : Reconstructed/Predicted data via
113         → given # of PCs
114
115     """
116
117     pca_mean = trained_pca.mean_
118     mean_removed = data - pca_mean
119     pca_components = trained_pca.components_[:number_PCs]
120
121     return mean_removed @ pca_components.T @ pca_components + pca_mean
122
123
124 def plot_faces(faces:np.ndarray,
125                 suptitle:str) -> None:
126     """
127
128     Given the face and its suptitle, plots the 6x6 grid.
129
130     Arguments:
131         - faces      (np.ndarray) : Face data to be plotted
132         - suptitle   (str)       : Suptitle of the visualization
133
134     Returns:
135         - None
136
137     """
138
139     fig, axes = plt.subplots(6, 6,
140                           figsize=(10,10),
141                           facecolor='white',
142                           subplot_kw= {
143                               'xticks': [],
144                               'yticks': []
145                           })
146
147     fig.suptitle(suptitle,
148                  fontsize = '14')
149
150     fig.tight_layout(rect = [0, 0, 1, .95])
151
152     for i, ax in enumerate(axes.flat):
153         ax.imshow(faces[i].reshape(32, 32).T, cmap='gray')
154         ax.set_xlabel(i+1)
155
156
157 # In[11]:

```

```

158
159
160 faces_PCA_10 = pca_reconstruction(faces,pca,10)
161 faces_PCA_25 = pca_reconstruction(faces,pca,25)
162 faces_PCA_50 = pca_reconstruction(faces,pca,50)
163 faces_PCA_100 = pca.inverse_transform(principalComponents)
164
165
166 # In[12]:
167
168
169 plot_faces(faces,suptitle = 'Original Versions of the First 36 Images')
170
171
172 # In[13]:
173
174
175 plot_faces(faces_PCA_10,suptitle = 'Reconstructed 36 images based on first 10
176   → PCs')
177
178 # In[14]:
179
180
181 plot_faces(faces_PCA_25,suptitle = 'Reconstructed 36 images based on first 25
182   → PCs')
183
184 # In[15]:
185
186
187 plot_faces(faces_PCA_50,suptitle = 'Reconstructed 36 images based on first 50
188   → PCs')
189
190 # In[16]:
191
192
193 plot_faces(faces_PCA_100, suptitle = 'Reconstructed 36 images based on first 100
194   → PCs')
195
196 # In[17]:
197
198
199 def squared_error(y_true:np.ndarray,y_pred:np.ndarray) -> np.ndarray:
200     """
201         Given the ground truth matrix and prediction, computes element wise
202         squared error.
203
204             Arguments:
205                 - y_true (np.ndarray) : ground truth
206                 - y_pred (np.ndarray) : prediction

```

```

207
208     Returns:
209         square_error (np.ndarray) : Point-wise MSE loss
210
211     """
212     assert y_true.shape == y_pred.shape, f'Mismatch Dimension!, {y_true.shape}'
213     ↪ does not match with {y_pred.shape}'
214     return (y_true - y_pred) ** 2
215
216 # In[18]:
217
218
219 mse_10 = squared_error(y_true = faces, y_pred = faces_PCA_10)
220 mse_25 = squared_error(y_true = faces, y_pred = faces_PCA_25)
221 mse_50 = squared_error(y_true = faces, y_pred = faces_PCA_50)
222
223
224 std_mse_10 = mse_10.mean(-1).std()
225 std_mse_25 = mse_25.mean(-1).std()
226 std_mse_50 = mse_50.mean(-1).std()
227
228 mean_mse_10 = mse_10.mean()
229 mean_mse_25 = mse_25.mean()
230 mean_mse_50 = mse_50.mean()
231
232
233 print(f'PCA reconstruction loss stats based on first 10 PCs, \n (mean,std) =
234     ↪ {mean_mse_10, std_mse_10}')
235 print(f'PCA reconstruction loss stats based on first 25 PCs, \n (mean,std) =
236     ↪ {mean_mse_25, std_mse_25}')
237 print(f'PCA reconstruction loss stats based on first 50 PCs, \n (mean,std) =
238     ↪ {mean_mse_50, std_mse_50}')
239
240 # ### Part C
241
242
243 fastIca_10 = FastICA(n_components = 10, random_state = 5)
244 fastIca_components_10 = fastIca_10.fit_transform(faces)
245
246
247 fig, axes = plt.subplots(2, 5, figsize=(10,4))
248 for i, ax in enumerate(axes.flat):
249     ax.imshow(fastIca_10.components_[i].reshape(32, 32).T, cmap = 'gray')
250     ax.axis('off')
251
252
253 # In[20]:
254
255
256 fastIca_25 = FastICA(n_components = 25, whiten = True, random_state = 5)

```

```

257 fastIca_components_25 = fastIca_25.fit_transform(faces)
258
259
260 fig, axes = plt.subplots(5, 5, figsize=(10,10))
261 for i, ax in enumerate(axes.flat):
262     ax.imshow(fastIca_25.components_[i].reshape(32, 32).T, cmap = 'gray')
263     ax.axis('off')
264
265
266 # In[21]:
267
268
269 fastIca_50 = FastICA(n_components = 50, random_state = 5)
270 fastIca_components_50 = fastIca_50.fit_transform(faces)
271
272
273 fig, axes = plt.subplots(5, 10, figsize=(20,10))
274 for i, ax in enumerate(axes.flat):
275     ax.imshow(fastIca_50.components_[i].reshape(32, 32).T, cmap = 'gray')
276     ax.axis('off')
277
278
279 # In[22]:
280
281
282 faces_ICA_10 = fastIca_10.inverse_transform(fastIca_components_10)
283 faces_ICA_25 = fastIca_25.inverse_transform(fastIca_components_25)
284 faces_ICA_50 = fastIca_50.inverse_transform(fastIca_components_50)
285
286
287 # In[23]:
288
289
290 plot_faces(faces_ICA_10, suptitle = 'FastICA reconstruction based on 10
291   ↪ independent components')
292
293
294 # In[24]:
295
296
297 plot_faces(faces_ICA_25, suptitle = 'FastICA reconstruction based on 25
298   ↪ independent components')
299
300
301
302 plot_faces(faces_ICA_50, suptitle = 'FastICA reconstruction based on 50
303   ↪ independent components')
304
305
306 # In[26]:
307
```

```

308 mse_10 = squared_error(y_true = faces, y_pred = faces_ICA_10)
309 mse_25 = squared_error(y_true = faces, y_pred = faces_ICA_25)
310 mse_50 = squared_error(y_true = faces, y_pred = faces_ICA_50)
311
312
313 std_mse_10 = mse_10.mean(-1).std()
314 std_mse_25 = mse_25.mean(-1).std()
315 std_mse_50 = mse_50.mean(-1).std()
316
317 mean_mse_10 = mse_10.mean()
318 mean_mse_25 = mse_25.mean()
319 mean_mse_50 = mse_50.mean()
320
321
322 print(f'ICA reconstruction loss stats based on first 10 ICs, \n (mean,std) =
    ↪ {mean_mse_10, std_mse_10}')
323 print(f'ICA reconstruction loss stats based on first 25 ICs, \n (mean,std) =
    ↪ {mean_mse_25, std_mse_25}')
324 print(f'ICA reconstruction loss stats based on first 50 ICs, \n (mean,std) =
    ↪ {mean_mse_50, std_mse_50}')
325
326
327 # In[27]:
328
329 """
330 """
331 fastICA can be seen as whitening (which can be achieved by PCA) plus an
    ↪ orthogonal rotation (an orthogonal rotation such that the estimated sources
    ↪ are as non-gaussian as possible).
332
333 The orthogonal rotation does not affect the reconstruction error of the ICA
    ↪ solution and hence you have the same reconstruction error for PCA and ICA.
334 """
335 """
336
337
338 # In[28]:
339
340
341 from sklearn.preprocessing import MinMaxScaler
342
343
344 # #### Part D
345
346 # In[29]:
347
348
349 faces = h5py.File('hw4_data1.mat','r')['faces'][:,].T
350 max_iter = 500
351 faces += np.abs(np.min(faces))
352 NMF_10 = NMF(n_components = 10, max_iter = max_iter)
353 NMF_components_10 = NMF_10.fit_transform(faces)
354
355

```

```

356 fig, axes = plt.subplots(2, 5, figsize=(10,4))
357 for i, ax in enumerate(axes.flat):
358     ax.imshow(NMF_10.components_[i].reshape(32, 32).T, cmap = 'gray')
359     ax.axis('off')
360
361
362 # In[30]:
363
364
365 NMF_25 = NMF(n_components = 25, max_iter = max_iter)
366 NMF_components_25 = NMF_25.fit_transform(faces)
367
368
369 fig, axes = plt.subplots(5, 5, figsize=(10,10))
370 for i, ax in enumerate(axes.flat):
371     ax.imshow(NMF_25.components_[i].reshape(32, 32).T, cmap = 'gray')
372     ax.axis('off')
373
374
375 # In[31]:
376
377
378 NMF_50 = NMF(n_components = 50, max_iter = max_iter)
379 NMF_components_50 = NMF_50.fit_transform(faces)
380
381
382 fig, axes = plt.subplots(5, 10, figsize=(12,6))
383 for i, ax in enumerate(axes.flat):
384     ax.imshow(NMF_50.components_[i].reshape(32, 32).T, cmap = 'gray')
385     ax.axis('off')
386
387
388 # In[32]:
389
390
391 faces = h5py.File('hw4_data1.mat','r')['faces'][:,].T
392
393 faces_NNMF_10 = NMF_10.inverse_transform(NMF_components_10) -
394     np.abs(np.min(faces))
395 faces_NNMF_25 = NMF_25.inverse_transform(NMF_components_25) -
396     np.abs(np.min(faces))
397 faces_NNMF_50 = NMF_50.inverse_transform(NMF_components_50) -
398     np.abs(np.min(faces))
399
400
401
402
403
404 # In[ ]:
405
406
407 # In[33]:

```

```

407 plot_faces(faces_NNMF_10, suptitle = 'NNMF Reconstruction of faces based on
408   ↪ 10MFs')
409
410 # In[34]:
411
412
413 plot_faces(faces_NNMF_25, suptitle = 'NNMF Reconstruction of faces based on
414   ↪ 25MFs')
415
416 # In[35]:
417
418
419 plot_faces(faces_NNMF_50, suptitle = 'NNMF Reconstruction of faces based on
420   ↪ 50MFs')
421
422 # In[36]:
423
424
425 mse_10 = squared_error(y_true = faces, y_pred = faces_NNMF_10)
426 mse_25 = squared_error(y_true = faces, y_pred = faces_NNMF_25)
427 mse_50 = squared_error(y_true = faces, y_pred = faces_NNMF_50)
428
429
430 std_mse_10 = mse_10.mean(-1).std()
431 std_mse_25 = mse_25.mean(-1).std()
432 std_mse_50 = mse_50.mean(-1).std()
433
434 mean_mse_10 = mse_10.mean()
435 mean_mse_25 = mse_25.mean()
436 mean_mse_50 = mse_50.mean()
437
438
439 print(f'NNMF reconstruction loss stats based on first 10 MFs, \n (mean,std) =
440   ↪ {mean_mse_10, std_mse_10}')
441 print(f'NNMF reconstruction loss stats based on first 25 MFs, \n (mean,std) =
442   ↪ {mean_mse_25, std_mse_25}')
443 print(f'NNMF reconstruction loss stats based on first 50 MFs, \n (mean,std) =
444   ↪ {mean_mse_50, std_mse_50}')
445
446 # #### Q2
447
448 # ### Part A
449
450 # In[38]:
451 def gauss_tuning(x:np.ndarray = np.linspace(-15, 16, 500),
452                   mu:float = 1,
453                   sigma:float = 1,
454                   A:float = 1) -> np.float16:

```

```

455 """
456     Gaussian shaped tuning function of a population of neurons.
457
458     Arguments:
459         x      (np.ndarray)    : The input stimulus parameters
460         A      (float)        : Gain of the Gaussian-shaped tuning curve
461         mu     (float)        : Mean of the Gaussian-shaped tuning curve
462         sigma  (float)       : Standard deviation of the Gaussian-shaped
463             → tuning curve
464
465     Returns:
466         response : Resulting neural response
467 """
468
469
470 # In[39]:
471
472
473 neural_responses = []
474 legends = []
475 stimuli = np.linspace(-15, 16, 500)
476 means = np.arange(-10, 11)
477 plt.figure(figsize=(10,5))
478
479 for mean in means:
480     response = gauss_tuning(mu = mean)
481     plt.plot(stimuli, response)
482     legends.append(f" Response with mean {mean}")
483
484 # Let's keep neural responses for future use:
485 neural_responses.append(response)
486
487 # Plot the tuning profiles
488 plt.plot(stimuli, np.mean(neural_responses, axis=0), color = '0')
489 legends.append(f" Average Response")
490 plt.legend(legends, loc="right", bbox_to_anchor=(1.4, 0.5))
491 plt.xlabel('Stimulus')
492 plt.ylabel('Activity')
493 plt.title('Tuning Curves of a Population of Neurons')
494 plt.grid()
495 plt.show()
496
497
498 # In[40]:
499
500
501 kwargs = dict(
502             color = '0',
503             marker= 'o',
504             markerfacecolor='red'
505         )
506 plt.figure(figsize=(10,5))
507 plt.plot(means, gauss_tuning(-1, mu = means), **kwargs)
508 plt.xlabel('Preferred Stimulus')

```

```

509 plt.ylabel('Population Response')
510 plt.title('Population Response to the Stimulus x = -1 vs Preferred Stimuli of
511    → Neurons')
512 plt.grid()
513 plt.show()
514
515 # #### Part B
516
517 #
518
519 # In[51]:
520
521
522 def WTA_decoder(stimuli:np.ndarray, response:np.ndarray) -> np.float16:
523     """
524         Given a population response and stimuli of the
525         neurons, compute the winner-take-all decoder that
526         estimates the actual stimulus as the preferred
527         stimulus of the neuron with maximum response
528
529             Arguments:
530                 stimuli (np.ndarray): The preferred stimuli of the neurons
531                 response (np.ndarray): The neural responses
532             Returns:
533                 stimulus (np.float16): the estimated input stimulus that maximizes
534         → the response
535         """
536
537         response += np.random.normal(loc = 0, scale = 1/20, size = (21,))
538
539         return stimuli[np.argmax(response)]
540
541
542
543 n_trials = 200
544 stimuli_interval = np.linspace(-5,5, 500).tolist()
545
546 # To keep 'responses', 'stimuli', 'WTA_stimuli', 'Errors'
547 stimuli_response_ = []
548
549 for stimuli in random.sample(stimuli_interval, n_trials):
550     response = gauss_tuning(stimuli, mu = means)
551     WTA_stimuli = WTA_decoder(means, response)
552     stimuli_response_.append((response, stimuli, WTA_stimuli, np.abs(WTA_stimuli
553         → - stimuli)))
554
555 # Tuples are gathered:
556 stimuli_response = list(zip(*stimuli_response_))
557
558 # In[53]:
559
```

```

560
561 fig = plt.figure(figsize=(10,5))
562 error = np.array(stimuli_response[-1])
563 fig.suptitle(f'Error Stats: (mean,std) =
564     {round(error.mean(),2),round(error.std(),2)})')
565 plt.scatter(range(n_trials), stimuli_response[1], marker="o", color="r", s=30,
566     linewidths=1)
567 plt.scatter(range(n_trials), stimuli_response[2], marker="x", color="green",
568     s=30, linewidths=1)
569 plt.xlabel('# of Trials')
570 plt.ylabel('Stimulus')
571 plt.title('Actual and WTA Estimated Stimuli Across Trials')
572 plt.legend(['Actual', 'WTA Estimated'], loc='upper right')
573 plt.show()
574
575
576 # In[57]:
577
578 print('Error Statistics for Winner Take All Decoder')
579 print('Mean of errors in stimuli estimation:', error.mean().round(5))
580 print('Standard deviation of errors in stimuli estimation :',
581     error.std().round(5))
582
583 # ### Part C
584
585 # In[45]:
586
587 def MLE_decoder(stimuli_interval:np.ndarray = np.linspace(-5,5, 500).tolist(),
588                 response:np.ndarray = None) -> np.float16:
589     """
590         Given a population response and stimuli of the
591         neurons, compute the MLE decoder that
592         estimates the actual stimulus as the preferred
593         stimulus of the neuron with maximum response
594
595         Arguments:
596             stimuli (np.ndarray): The preferred stimuli of the neurons
597             response (np.ndarray): The neural responses
598         Returns:
599             stimulus (np.float16): the estimated input stimulus that maximizes
600             the response
601             """
602
603 logs = list()
604
605 for stim in stimuli_interval:
606     log = sum((response_ - gauss_tuning(stim, mu_)) ** 2 for response_, mu_
607     in zip(response, means))

```

```

608     logs.append(log)
609
610     idx_stim_max = np.argmin(logs)
611     return stimuli_interval[idx_stim_max]
612
613
614
615 # In[59]:
616
617
618 MLE_logs_ = list()
619 for neural_activity in stimuli_response_:
620     response, stimulus = neural_activity[:2]
621     estimated_stimulus_MLE = MLE_decoder(stimuli_interval, response)
622     MLE_logs_.append((stimulus, estimated_stimulus_MLE, np.abs(stimulus -
623         ↪ estimated_stimulus_MLE)))
624
625 # Tuples are gathered:
626 MLE_logs = list(zip(*MLE_logs_))
627
628
629 # In[60]:
630
631
632 fig = plt.figure(figsize=(10,5))
633 error = np.array(MLE_logs[2])
634 fig.suptitle(f'Error Stats: (mean,std) =
635     ↪ {round(error.mean(),2),round(error.std(),2)}')
636 plt.scatter(range(n_trials), MLE_logs[0], marker="o", color="r", s=30,
637     ↪ linewidths=1)
638 plt.scatter(range(n_trials), MLE_logs[1], marker="x", color="green", s=30,
639     ↪ linewidths=1)
640 plt.xlabel('# of Trials')
641 plt.ylabel('Stimulus')
642 plt.title('Actual and MLE Estimated Stimuli Across Trials')
643 plt.legend(['Actual', 'MLE Estimated'], loc='upper right')
644 plt.show()
645
646
647 # In[61]:
648
649
650 print('Error Statistics for MLE Decoder')
651 print('Mean of errors in stimuli estimation:', error.mean().round(5))
652 print('Standard deviation of errors in stimuli estimation :',
653     ↪ error.std().round(5))
654
655
656 # In[ ]:
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
999

```

```

657 # #### Part D
658 # In[62]:
659
660
661
662
663 def MAP_decoder(stimuli_interval:np.ndarray = np.linspace(-5,5, 500).tolist(),
664                 response:np.ndarray = None) -> np.float16:
665     """
666
667     Given a population response and stimuli of the
668     neurons, compute the MAP decoder that
669     estimates the actual stimulus as the preferred
670     stimulus of the neuron with maximum response
671
672     Arguments:
673         stimuli (np.ndarray): The preferred stimuli of the neurons
674         response (np.ndarray): The neural responses
675
676     Returns:
677         stimulus (np.float16): the estimated input stimulus that maximizes
678         the response
679     """
680
681     logs = list()
682
683     for stim in stimuli_interval:
684
685         log = sum((r - gauss_tuning(stim, m)) ** 2 for r, m in zip(response,
686                         means))
686         log = log * 200 + (stim ** 2) / 10
687
688         logs.append(log)
689
690     idx_stim_max = np.argmin(logs)
691     return stimuli_interval[idx_stim_max]
692
693
694
695 # In[63]:
696
697
698 MAP_logs_ = list()
699 for neural_activity in stimuli_response_:
700     response, stimulus = neural_activity[:2]
701     estimated_stimulus_MAP = MAP_decoder(stimuli_interval, response)
702     MAP_logs_.append((stimulus, estimated_stimulus_MAP, np.abs(stimulus -
703                         estimated_stimulus_MAP)))
704
705
706 # Tuples are gathered:
707 MAP_logs = list(zip(*MAP_logs_))
708

```

```

708 # In[64]:
710
711
712 fig = plt.figure(figsize=(10,5))
713 error = np.array(MAP_logs[2])
714 fig.suptitle(f'Error Stats: (mean,std) =
    ↪ {round(error.mean(),2),round(error.std(),2)}')
715 plt.scatter(range(n_trials), MAP_logs[0], marker="o", color="r", s=30,
    ↪ linewidths=1)
716 plt.scatter(range(n_trials), MAP_logs[1], marker="x", color="green", s=30,
    ↪ linewidths=1)
717 plt.xlabel('# of Trials')
718 plt.ylabel('Stimulus')
719 plt.title('Actual and MAP Estimated Stimuli Across Trials')
720 plt.legend(['Actual', 'MAP Estimated'], loc='upper right')
721 plt.show()
722
723
724 # In[66]:
725
726
727 print('Error Statistics for MAP Decoder')
728 print('Mean of errors in stimuli estimation:', error.mean().round(5))
729 print('Standard deviation of errors in stimuli estimation :',
    ↪ error.std().round(5))
730
731
732 # #### Part E
733
734 # In[144]:
735
736
737 stds = [0.1, 0.2, 0.5, 1.0, 2.0, 5.0]
738 bias_outer = []
739
740 for stimuli in random.sample(stimuli_interval.tolist(), n_trials):
    bias_inner = []
    for std in stds:
        response = gauss_tuning(stimuli,means,sigma=std)
        response += np.random.normal(loc = 0, scale = 1/20, size = (21,))
        stimulus_MLE = MLE_decoder(response=response)
        bias = np.abs(stimuli - stimulus_MLE)
        bias_inner.append(bias)
741
742
743     bias_outer.append(tuple(bias_inner))
744
745
746
747
748
749
750
751 # Tuples are gathered:
752 sigmas_response_bias = list(zip(*bias_outer))
753
754
755 # In[227]:
756
757

```

```

758 error_stats = {}
759
760 for bias, sigma in zip(sigmas_responses, stds):
761
762     ave_error = np.mean(bias).round(3)
763     std_error = np.std(bias).round(3)
764     error_stats[sigma] = ave_error, std_error
765
766 error_stats = pd.DataFrame(error_stats).T
767 error_stats.columns = ['Mean Error', 'STD Error']
768 error_stats.index = stds
769 error_stats.index.name = 'Sigma Values'
770 error_stats.sort_values('Mean Error', ascending = True, inplace = True)
771 error_stats.head(n = 6)
772
773
774 #
775
776 # In[230]:
777
778
779 fig, ax = plt.subplots(figsize = (10,5))
780 error_stats.plot.kde(ax=ax)
781 error_stats.plot.hist(density=True, ax = ax)
782 ax.set_ylabel('Error Stats')
783 ax.set_xlabel('Sigma values $\sigma$')
784 ax.grid()
785
786
787 fig, ax = plt.subplots(figsize = (10,5))
788 error_stats.plot(ax=ax,marker = 'o')
789 ax.set_ylabel('Error Stats')
790 ax.set_xlabel('Sigma values $\sigma$')
791 ax.grid()
792
793
794
795 # In[ ]:
```

REFERENCES

- [1] Wikipedia contributors. *Curse of dimensionality* — Wikipedia, The Free Encyclopedia. [Online; accessed 25-April-2021]. 2021. URL: https://en.wikipedia.org/w/index.php?title=Curse_of_dimensionality&oldid=1009164221.
- [2] Wikipedia contributors. *Independent component analysis* — Wikipedia, The Free Encyclopedia. [Online; accessed 25-April-2021]. 2021. URL: https://en.wikipedia.org/w/index.php?title=Independent_component_analysis&oldid=1007441084.
- [3] Wikipedia contributors. *Winner-take-all (computing)* — Wikipedia, The Free Encyclopedia. [Online; accessed 27-April-2021]. 2019. URL: [https://en.wikipedia.org/w/index.php?title=Winner-take-all_\(computing\)&oldid=913093990](https://en.wikipedia.org/w/index.php?title=Winner-take-all_(computing)&oldid=913093990).